



# AWS Serverless

## TV2 Workshop

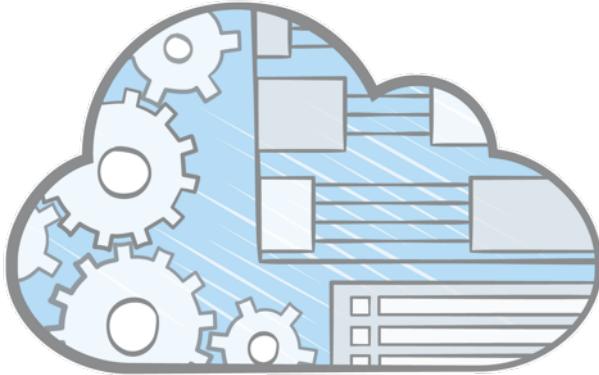
Tobias Börjeson, [borjeson@amazon.com](mailto:borjeson@amazon.com)

6 March, 2017

# Agenda

## Technical Workshop

- Intro
- AWS Technical Overview
- AWS Account Strategies
- Serverless Architectures
- CI/CD Pipelines
- Example Project

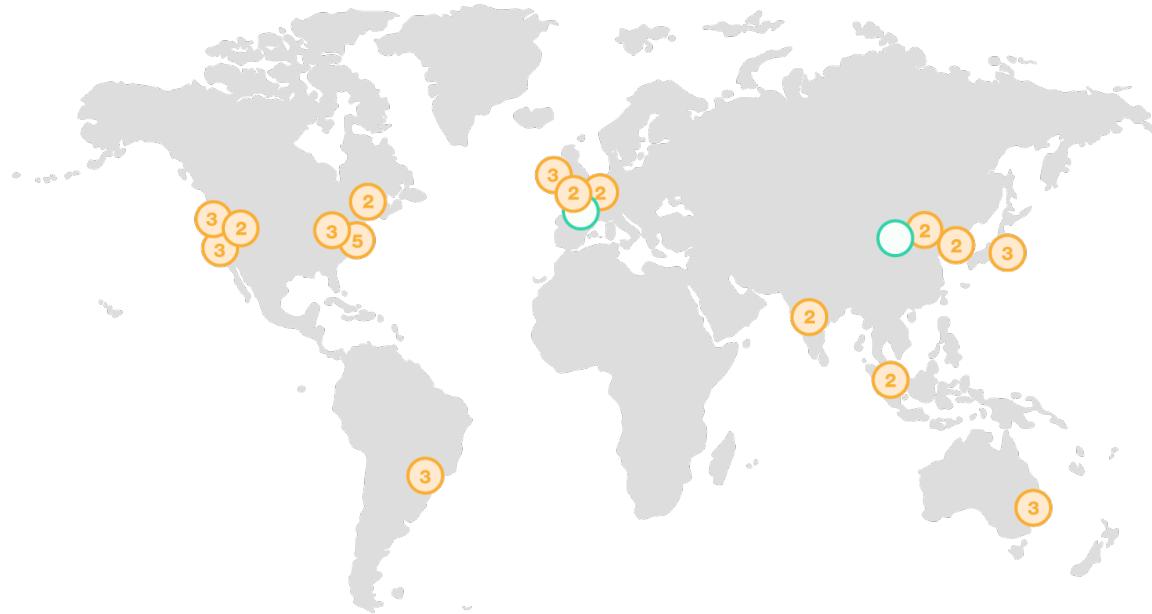


Objectives: Provide understanding of Serverless architectures on AWS, with focus on best practices and examples from a Java perspective

# AWS Technical Overview

# AWS Global Infrastructure

16 Regions – 42 Availability Zones – 68 Edge Locations



## Region & Number of Availability Zones

<b>AWS GovCloud (2)</b>	<b>EU</b>
	Ireland (3)
	Frankfurt (2)
<b>US West</b>	<b>London (2)</b>
Oregon (3)	
Northern California (3)	
	<b>Asia Pacific</b>
	Singapore (2)
<b>US East</b>	Sydney (2), Tokyo (3),
N. Virginia (5), Ohio (3)	Seoul (2), Mumbai (2)
<b>Canada</b>	<b>China</b>
Central (2)	Beijing (2)
<b>South America</b>	
São Paulo (3)	

## Announced Regions

Paris, Ningxia



# Europe / Middle East / Africa

## EU (Ireland) Region

EC2 Availability Zones: 3

## EU (Frankfurt) Region

EC2 Availability Zones: 2

## EU (London) Region

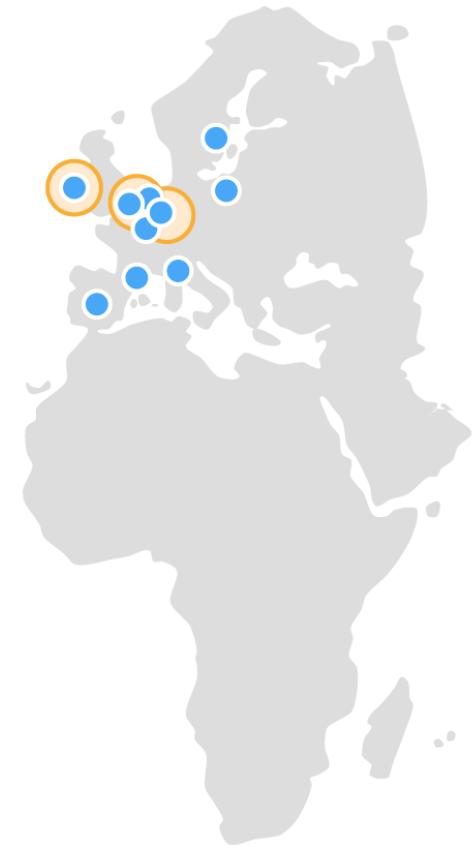
EC2 Availability Zones: 2

## EU (Paris) Region

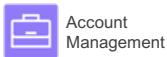
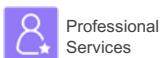
*Announced*

## AWS Edge Locations

Amsterdam, The Netherlands (2), Berlin, Germany, Dublin, Ireland, Frankfurt, Germany (5), London, England (4), Madrid, Spain, Marseille, France, Milan, Italy, Paris, France (2), Stockholm, Sweden, and Warsaw, Poland



TECHNICAL &  
BUSINESS  
SUPPORT



HYBRID  
ARCHITECTURE



ANALYTICS

- Data Warehousing
- Business Intelligence
- Hadoop/Spark
- Streaming Data Analysis
- Streaming Data Collection
- Machine Learning
- Elastic Search

APP SERVICES

- Queuing & Notifications
- Workflow
- Search
- Email
- Transcoding

MOBILE SERVICES

- API Gateway
- Identity
- Sync
- Mobile Analytics
- Single Integrated Console
- Push Notifications

DEVELOPMENT & OPERATIONS

- One-click App Deployment
- DevOps Resource Management
- Application Lifecycle Management
- Containers
- Triggers
- Resource Templates

IoT

- Rules Engine
- Device Shadows
- Device SDKs
- Device Gateway
- Registry

ENTERPRISE APPS

- Virtual Desktops
- Sharing & Collaboration
- Corporate Email
- Backup

SECURITY & COMPLIANCE



CORE SERVICES



INFRASTRUCTURE



# Spectrum of AWS offerings

“On EC2”



Amazon EC2



Microsoft SQL  
Server



Managed



Amazon  
EMR



Amazon  
Elasticsearch  
Service



Amazon  
ElastiCache



Amazon  
RDS



Amazon  
Redshift

Serverless



AWS  
Lambda



Amazon  
Cognito



Amazon  
Kinesis



Amazon  
S3



Amazon  
DynamoDB



Amazon  
SQS



Amazon  
API  
Gateway



Amazon  
CloudWatch



AWS IoT



# AWS building blocks

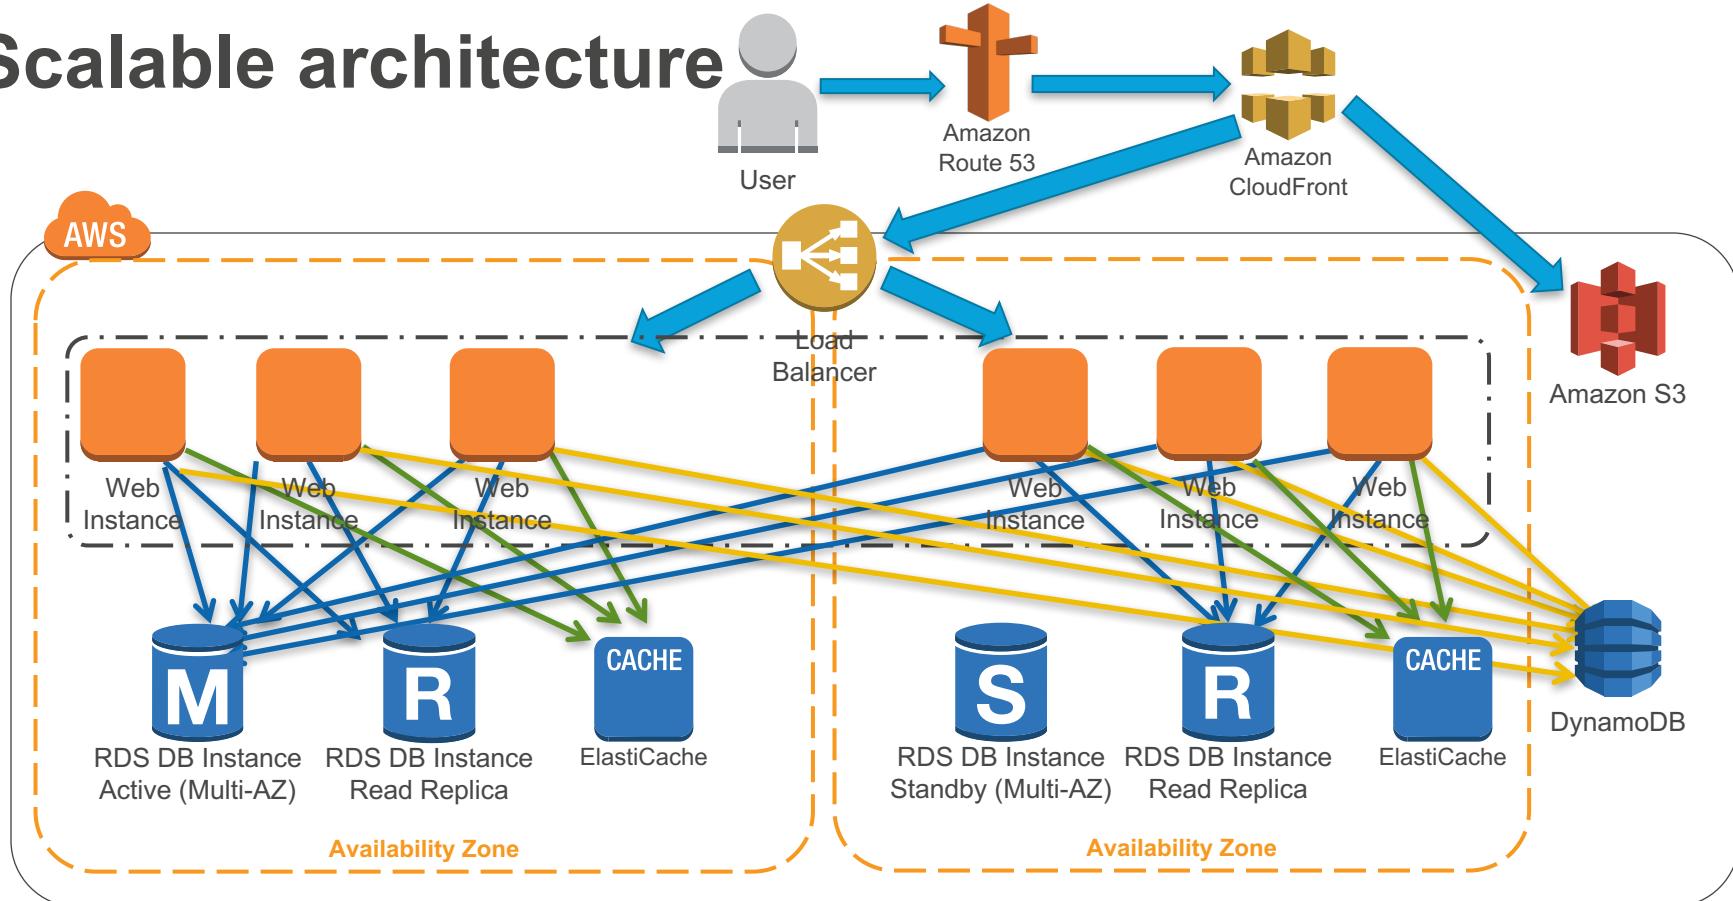
## Inherently highly available and fault-tolerant services

- ✓ Amazon CloudFront
- ✓ Amazon Route 53
- ✓ Amazon S3
- ✓ Amazon DynamoDB
- ✓ Elastic Load Balancing
- ✓ Amazon EFS
- ✓ AWS Lambda
- ✓ Amazon SQS
- ✓ Amazon SNS
- ✓ Amazon SES
- ✓ Amazon SWF
- ✓ ...

## Highly available with the right architecture

- ▶ Amazon EC2
- ▶ Amazon EBS
- ▶ Amazon RDS
- ▶ Amazon VPC

# Scalable architecture



# AWS application management solutions

Higher-level services



AWS  
Elastic Beanstalk



AWS  
OpsWorks

Do it yourself



AWS  
CloudFormation



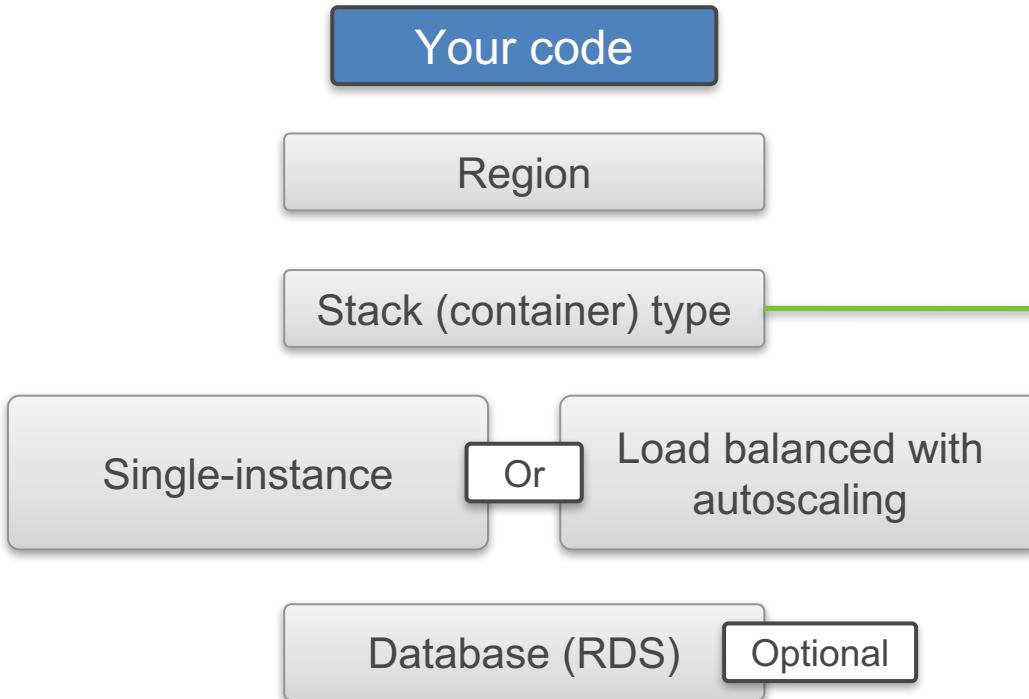
Amazon EC2

Convenience

Control

# AWS Elastic Beanstalk

- 01
- 02
- 03
- 04



# Loose coupling = winning

DON'T REINVENT THE WHEEL

- Email
- Queuing
- Transcoding
- Search
- Databases
- Monitoring
- Metrics
- Logging
- Compute



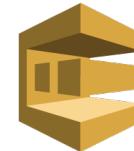
AWS Lambda



Amazon SNS



Amazon  
ElasticSearch



Amazon SQS



Amazon SES



Amazon SWF

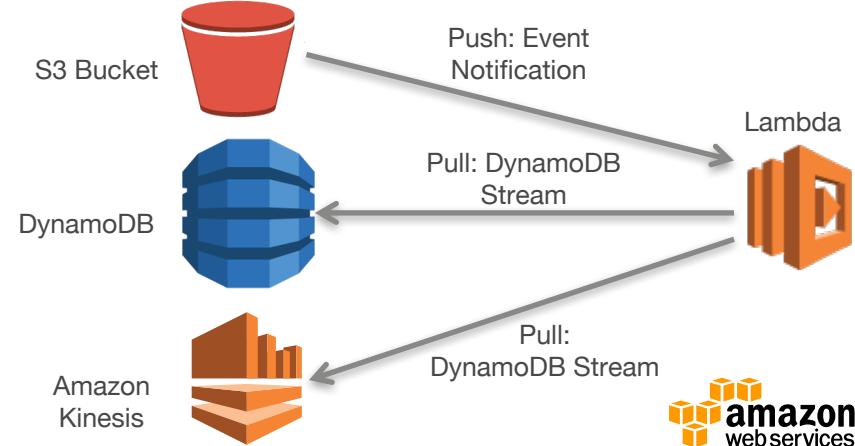
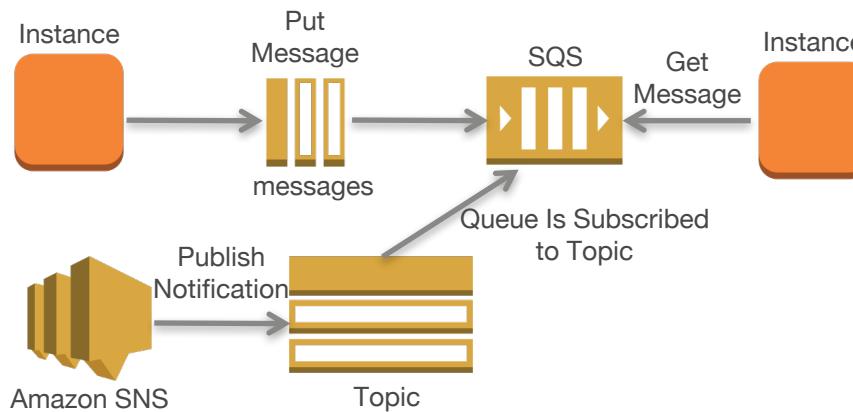


Amazon  
Elastic  
Transcoder

# Loose coupling sets you free!

The looser they're coupled, the bigger they scale

- Independent components
- Design everything as a black box
- Decouple interactions
- Favor services with built-in redundancy and scalability rather than building your own



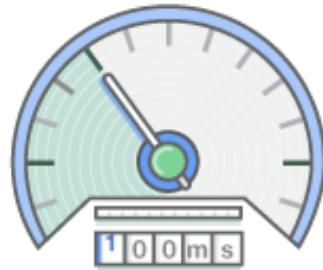
# Serverless patterns built with functions

Functions are the unit of deployment and scale

Scales per request - users cannot over or under-provision

Never pay for idle

Skip the boring parts; skip the hard parts



# Working with AWS Lambda

## EVENT SOURCE



Changes in  
data state



Requests to  
endpoints



Changes in  
resource state



## FUNCTION

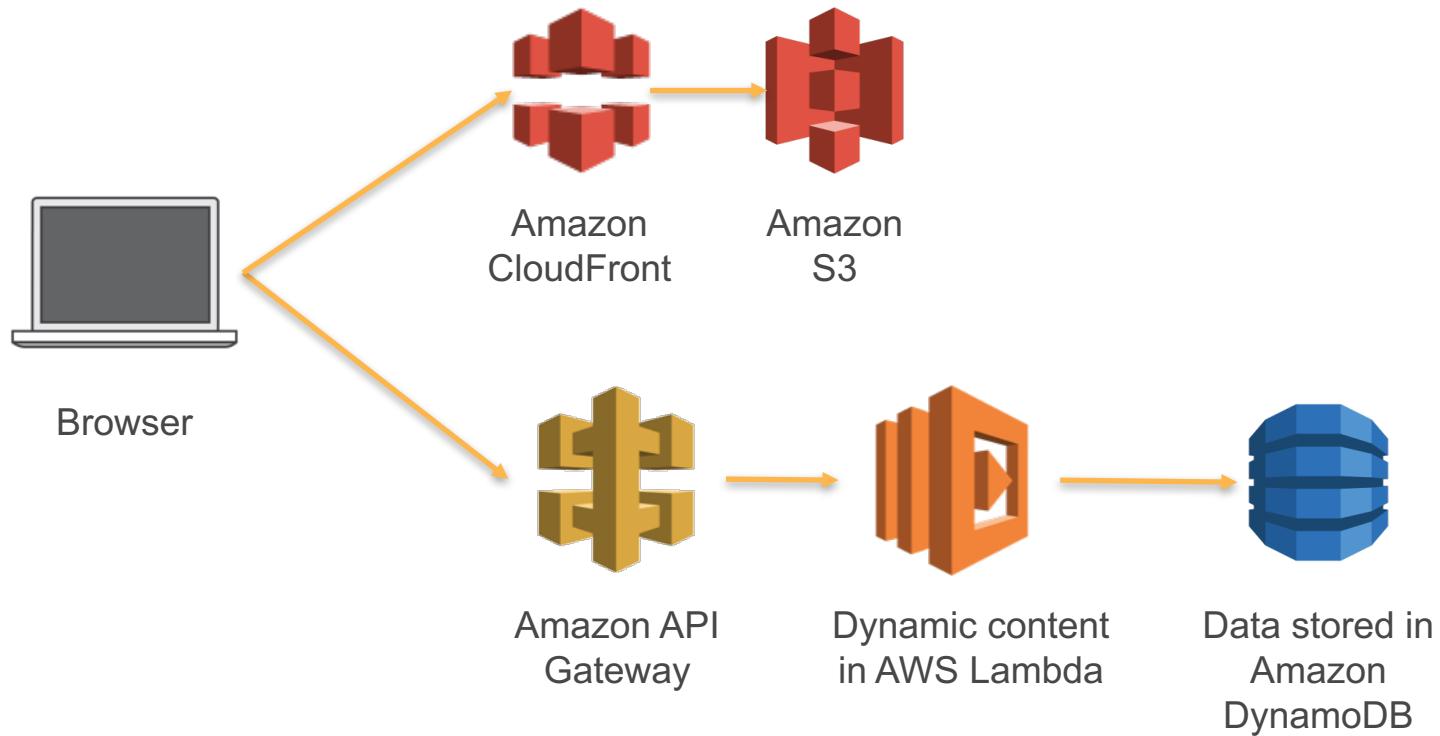


Node  
Python  
Java  
C#

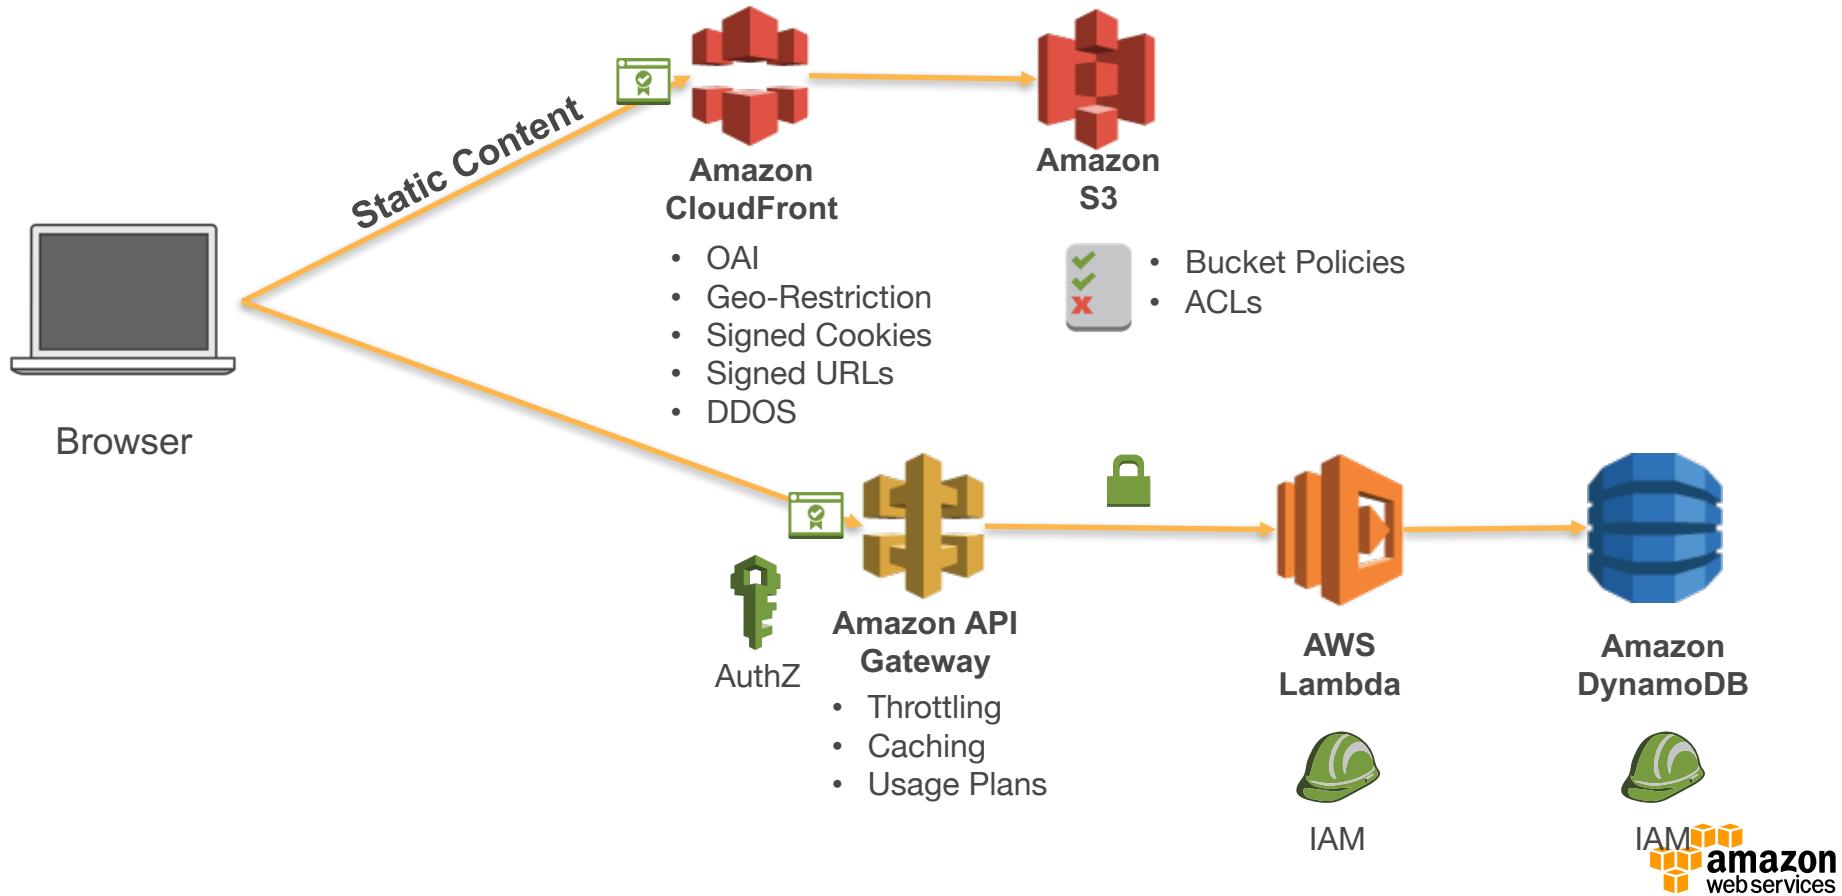
## SERVICES (ANYTHING)



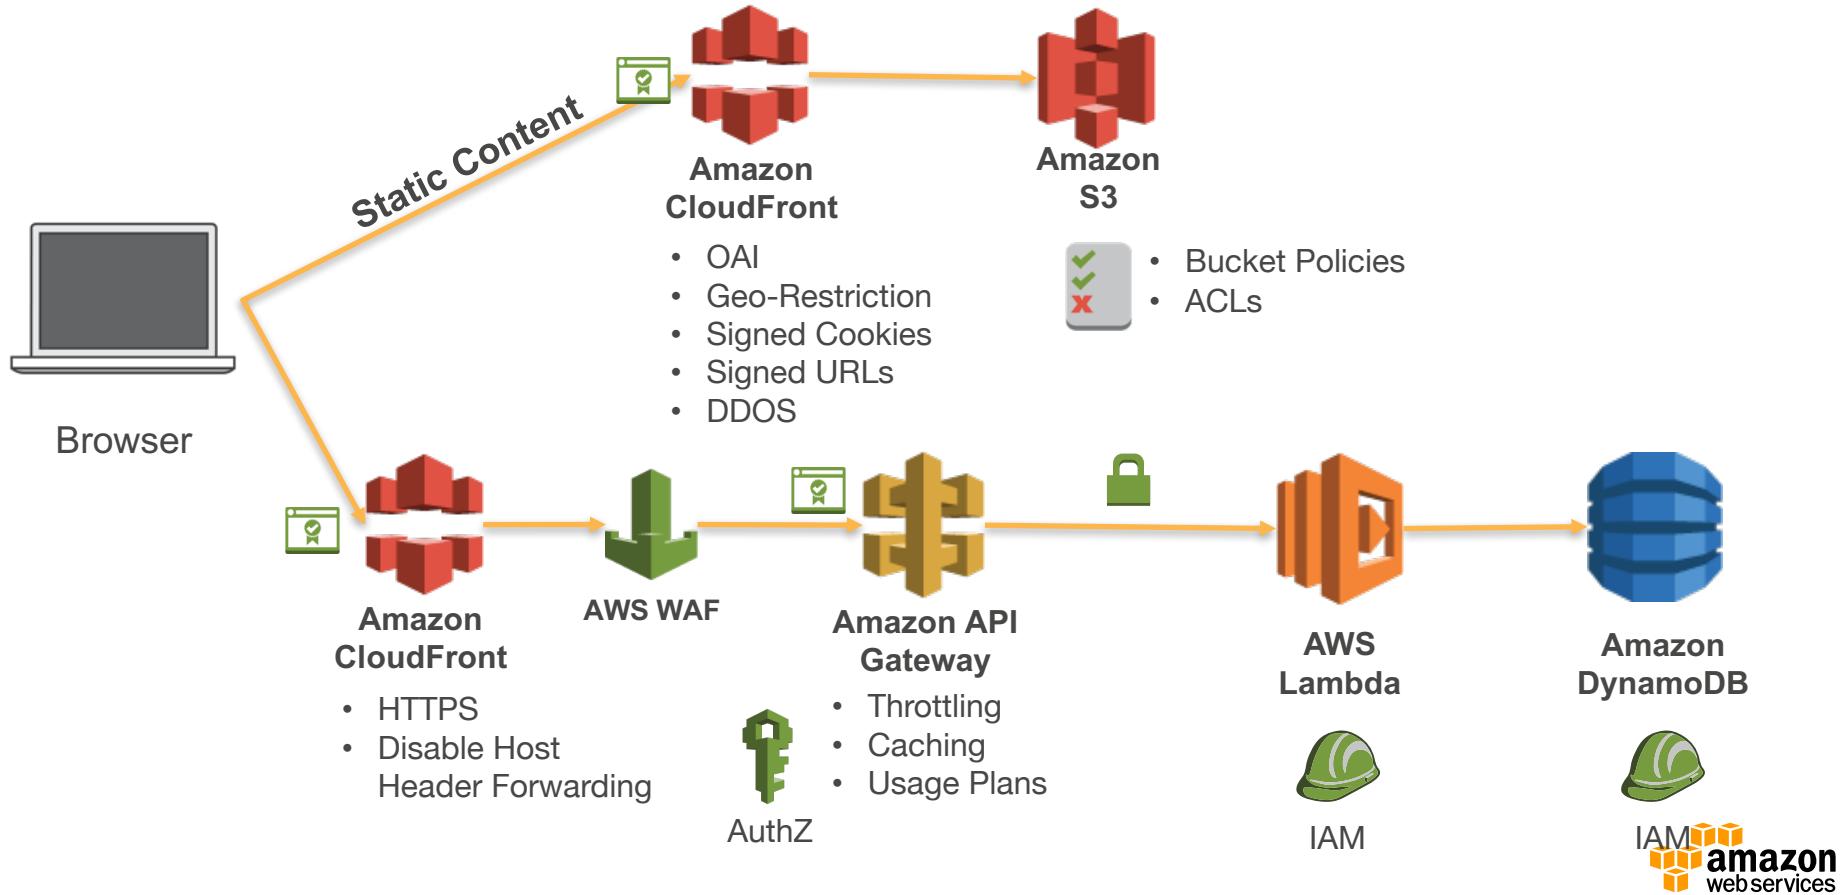
# Serverless Web application



# Serverless web app security

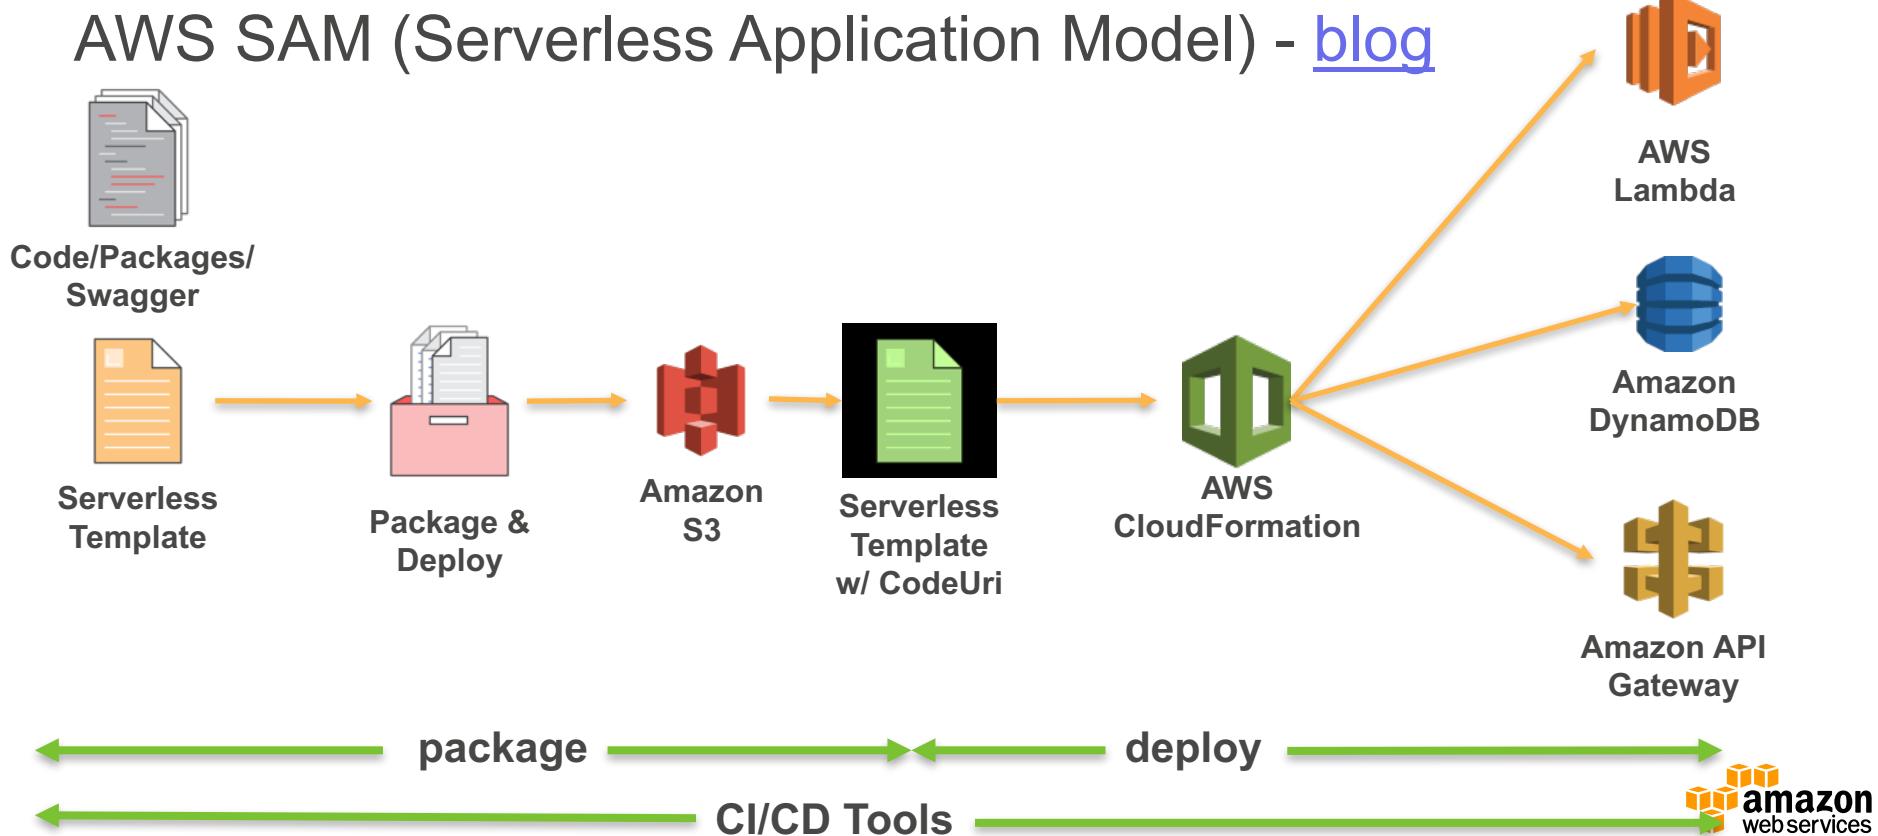


# Serverless web app security



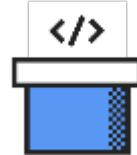
# Serverless web app lifecycle management

AWS SAM (Serverless Application Model) - [blog](#)



# AWS Code\* services

Software release steps:



AWS CodePipeline

Commit

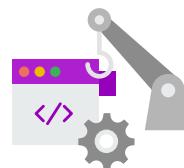
Build

Test

Production



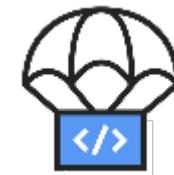
AWS CodeCommit



AWS CodeBuild

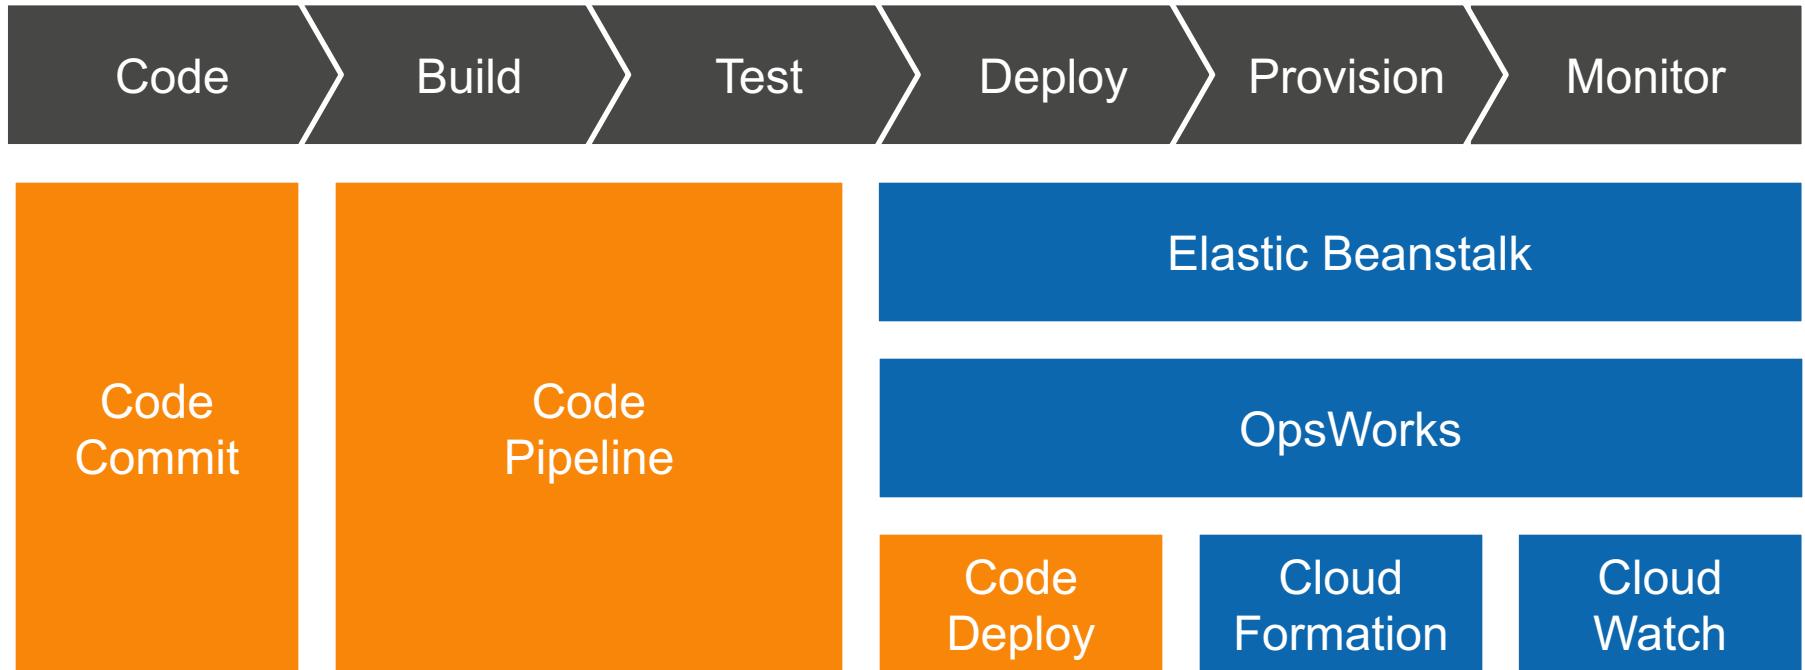


Third Party  
Tooling



AWS CodeDeploy

# AWS Code, Deployment & Management Services



# AWS Account Setup

# Free Tier

- Includes most of the AWS services
- Available for all new accounts
- Good for one year from the day the account is created
- **Everything we show today can be done within the Free Tier!**
- More details at <http://aws.amazon.com/free>

# Step 1: Sign up and configure basic security

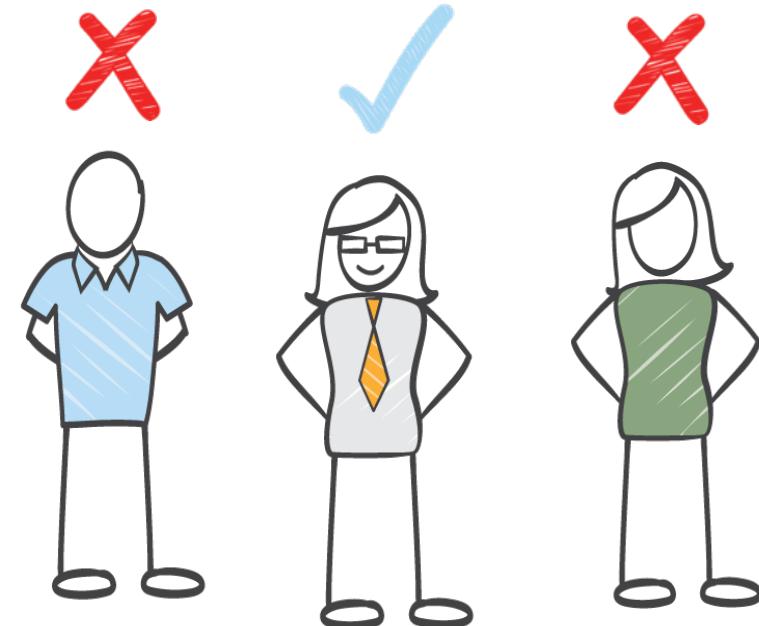
- Sign up though <https://aws.amazon.com>
- You need a credit card
- There will be a phone verification



# Basic security: Creating IAM users

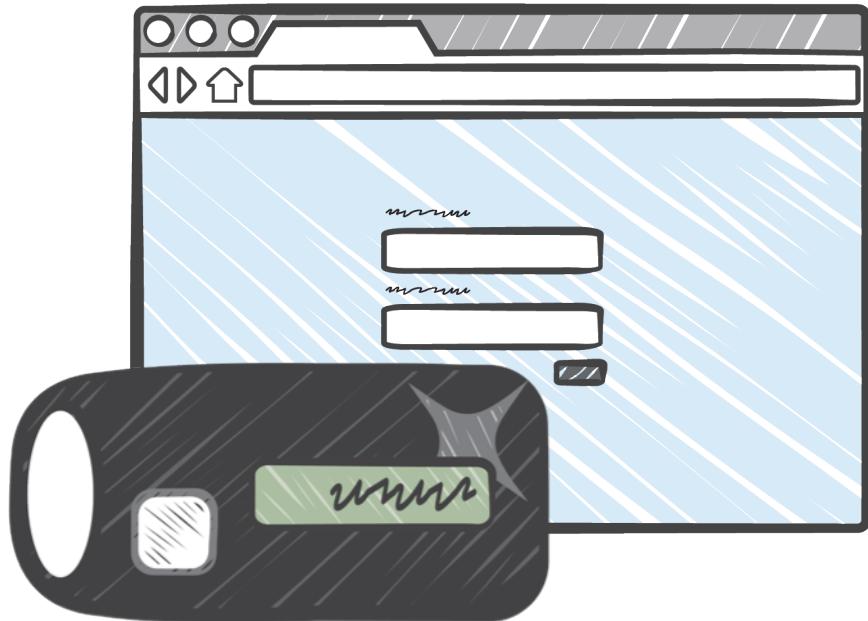
Using AWS Identity and Access Management (IAM), you can create and manage AWS users and groups.

You can control what resources each user has access to so you can avoid overly permissive accounts.



# Enabling MFA

AWS allows you to require **Multi-Factor Authentication** (MFA) for your users through physical-based or software-based single use login tokens to thwart stolen passwords and key loggers as an attack vector.



# AWS Account Strategies

# Characteristics of an AWS account



Security boundary



Resource isolation



Billing separation

# Account Structure Considerations

## IAM users

- Across single or multiple AWS accounts?

## AD Federation

- To federate or not to federate?

## Tool Chain

- Shared across all environments?

# Approaches to account structure



One  
Account



-



1000's of  
Accounts



# Why one isn't enough



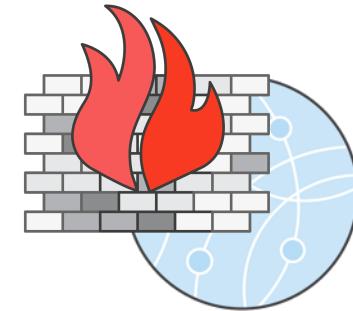
Many Teams



Security Controls



Billing



Isolation



Business Process

# Multiple accounts

## Pros

- + Complete security and resources isolation
- + Smaller blast radius
- + Simplified billing per account

## Cons

- Setup and operation overhead
- More complex security policies across accounts

# Goals of a multi-account architecture

- Automated setup of accounts
- Scalable
- Self-service
- Guardrails while enabling innovation
- Auditable for security and compliance
- Flexible

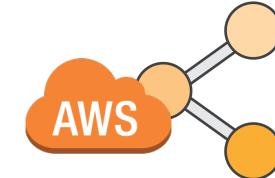
# What accounts should I create?



Billing



Security



Shared Services



Sandbox



Non-prod

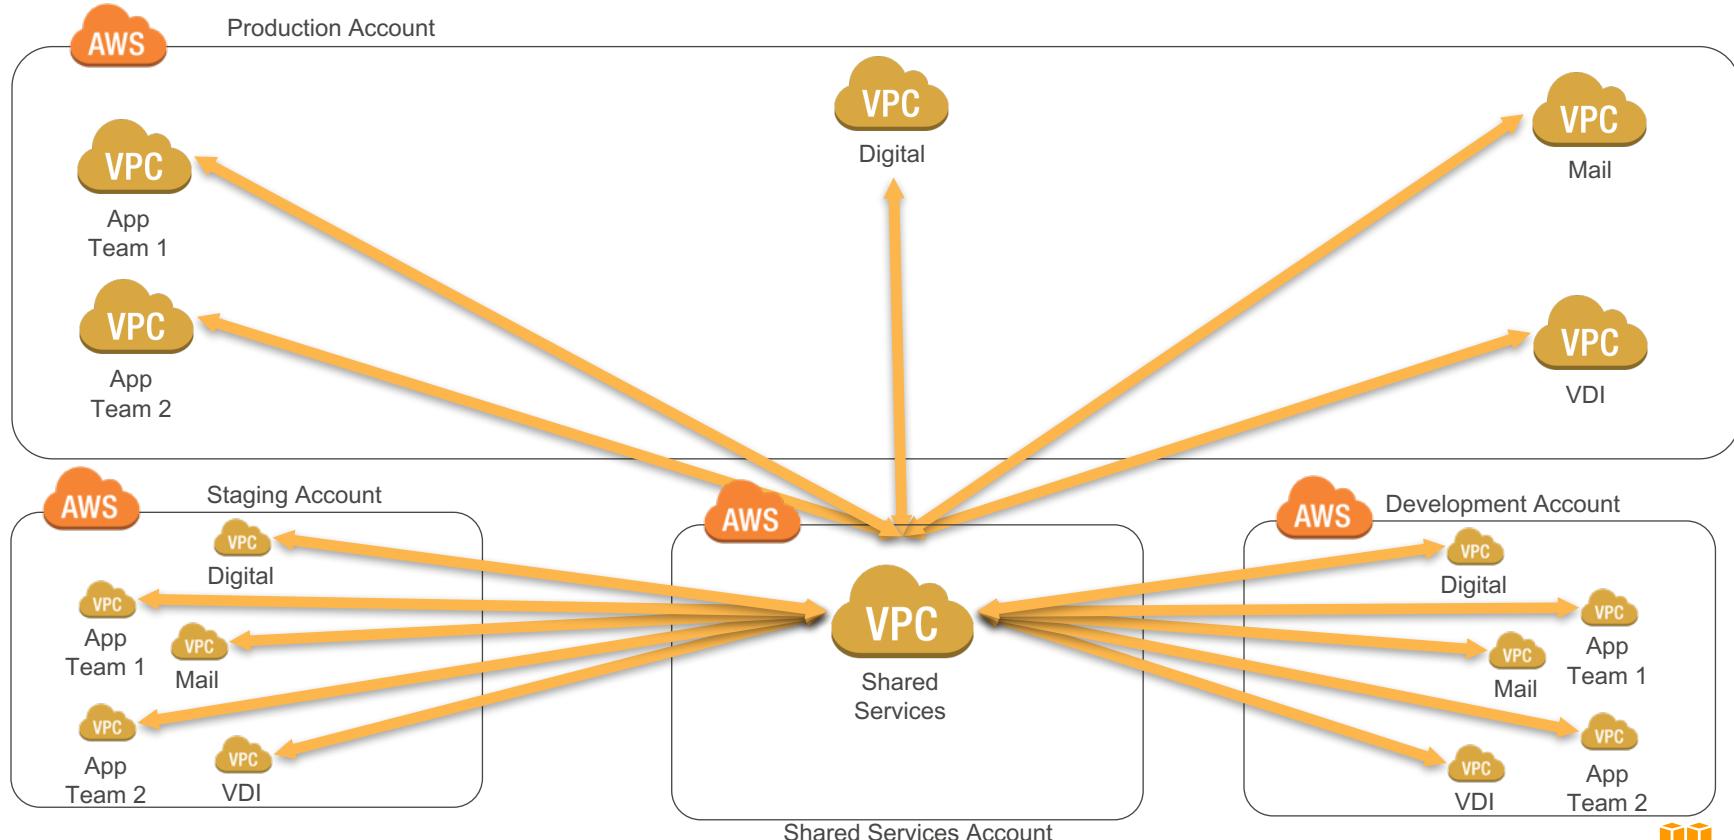


Prod

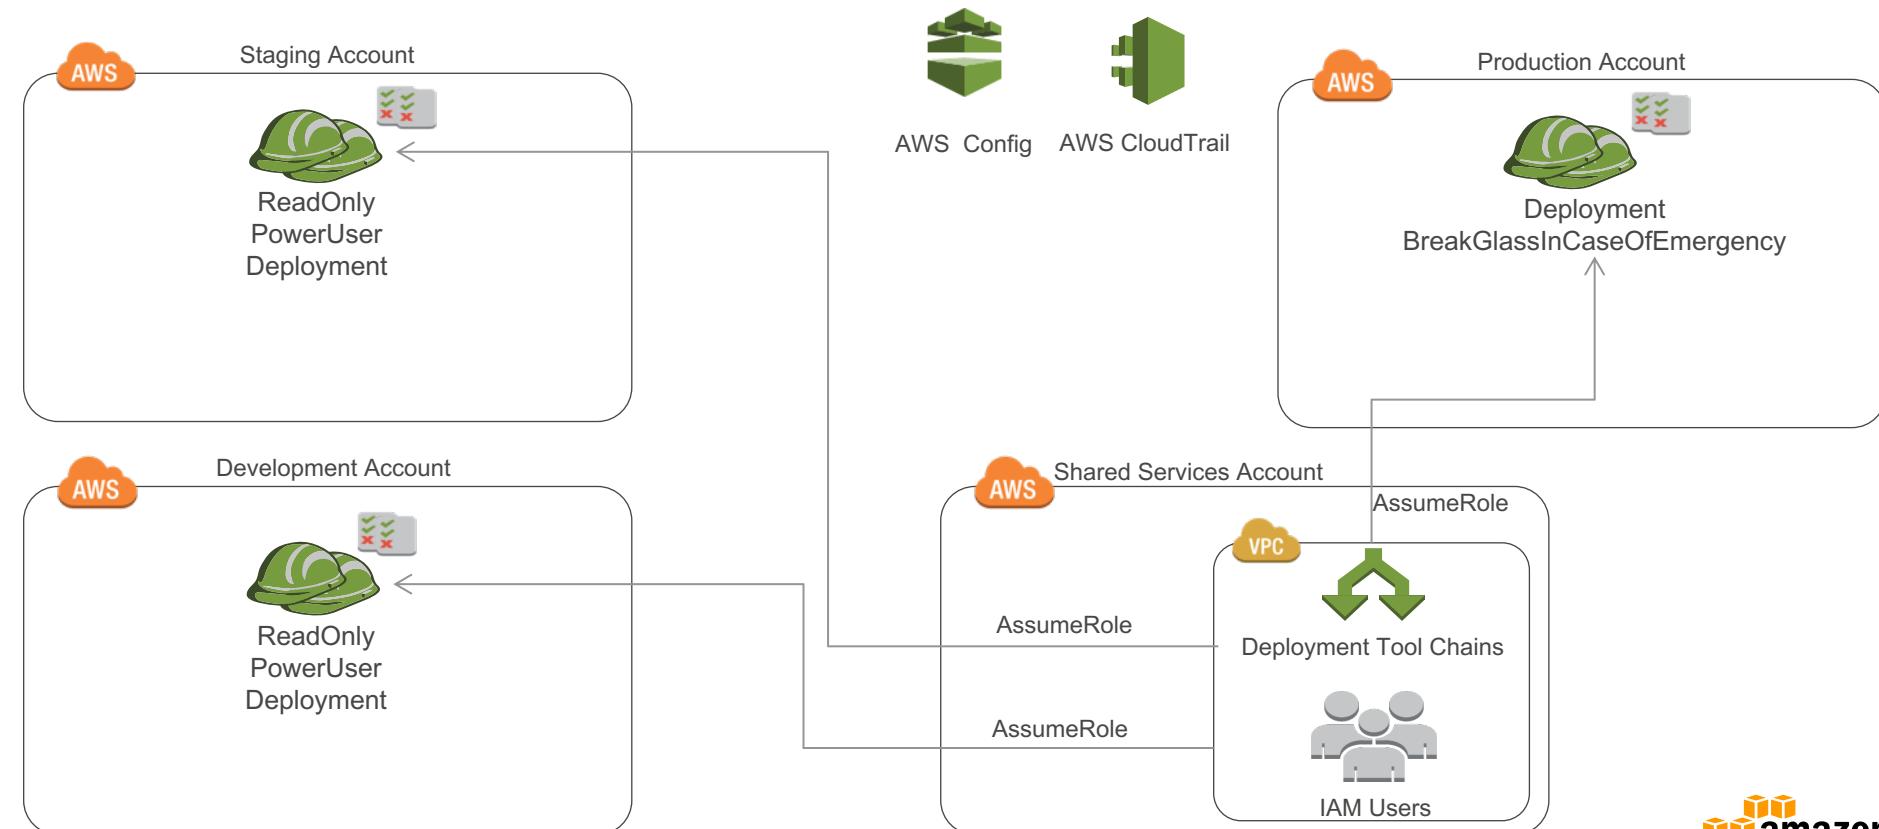


Other

# Central Account Structure



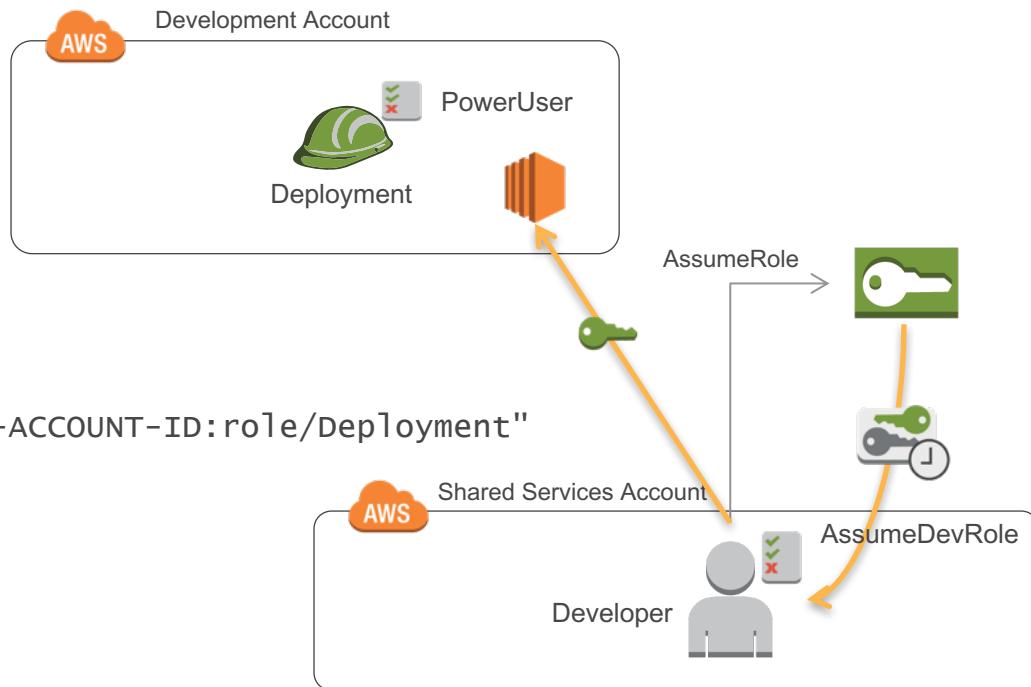
# Central User Management



# Grant User Access



```
"AssumeDeveloperRolePolicy": {  
    "Properties": {  
        "PolicyName": "AssumeDevRole",  
        "Users": [ "Developer" ],  
        "PolicyDocument": {  
            "Version": "2012-10-17",  
            "Statement": [  
                {  
                    "Effect": "Allow",  
                    "Action": "sts:AssumeRole"  
                    "Resource": "arn:aws:iam::DEV-ACCOUNT-ID:role/deployment"  
                }  
            ]  
        }  
    },  
    "Type": "AWS::IAM::Policy"  
}
```



# Demo AWS Console and CLI

# Additional Resources

- [AWS Getting Started](#)
- [AWS Multiple Account Billing Strategy](#)
- [Set up AWS CLI](#)

# Building Serverless APIs

# Building blocks for serverless applications

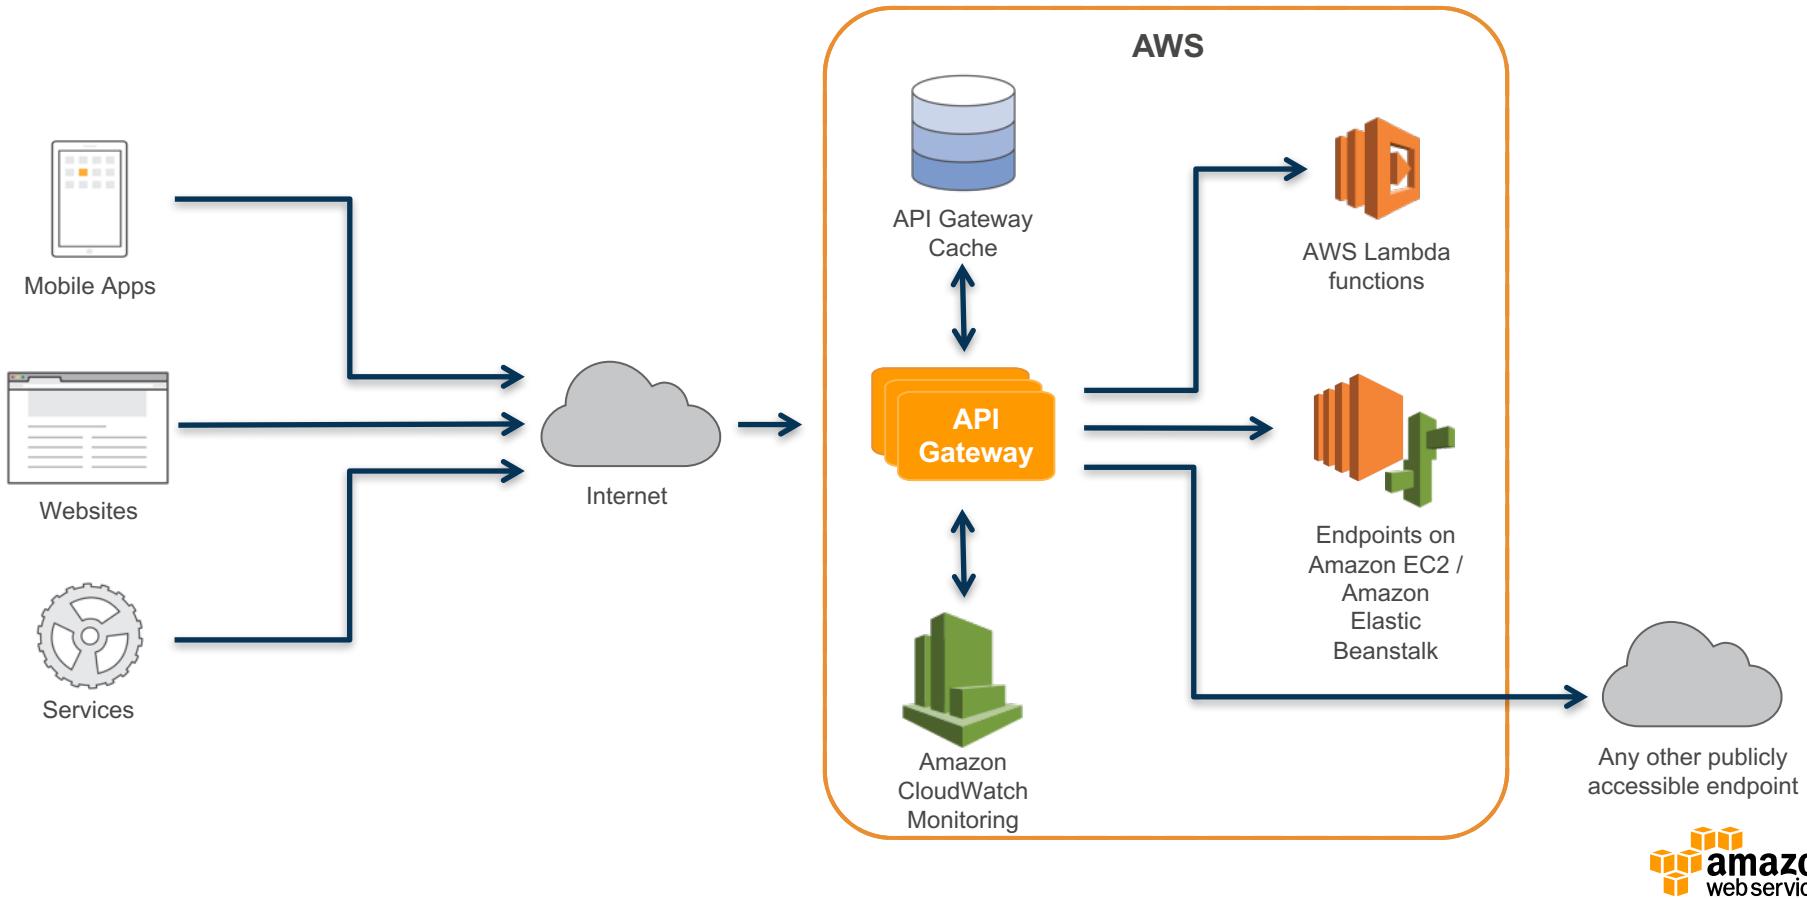
Compute	Storage	Database
 AWS Lambda	 Amazon S3	 Amazon DynamoDB
API Proxy	Messaging and Queues	Analytics
 Amazon API Gateway	 Amazon SQS  Amazon SNS	 Amazon Kinesis
Orchestration and State Management	Monitoring and Debugging	
 AWS Step Functions	 AWS X-Ray	

# Serverless APIs Fundamentals



# API Gateway

# API Gateway: Serverless APIs



# Introducing Amazon API Gateway



-  Host multiple versions and stages of your APIs
-  Create and distribute API Keys to developers
-  Leverage AWS Sigv4 to authorize access to APIs
-  Throttle and monitor requests to protect your backend
-  Utilizes AWS Lambda

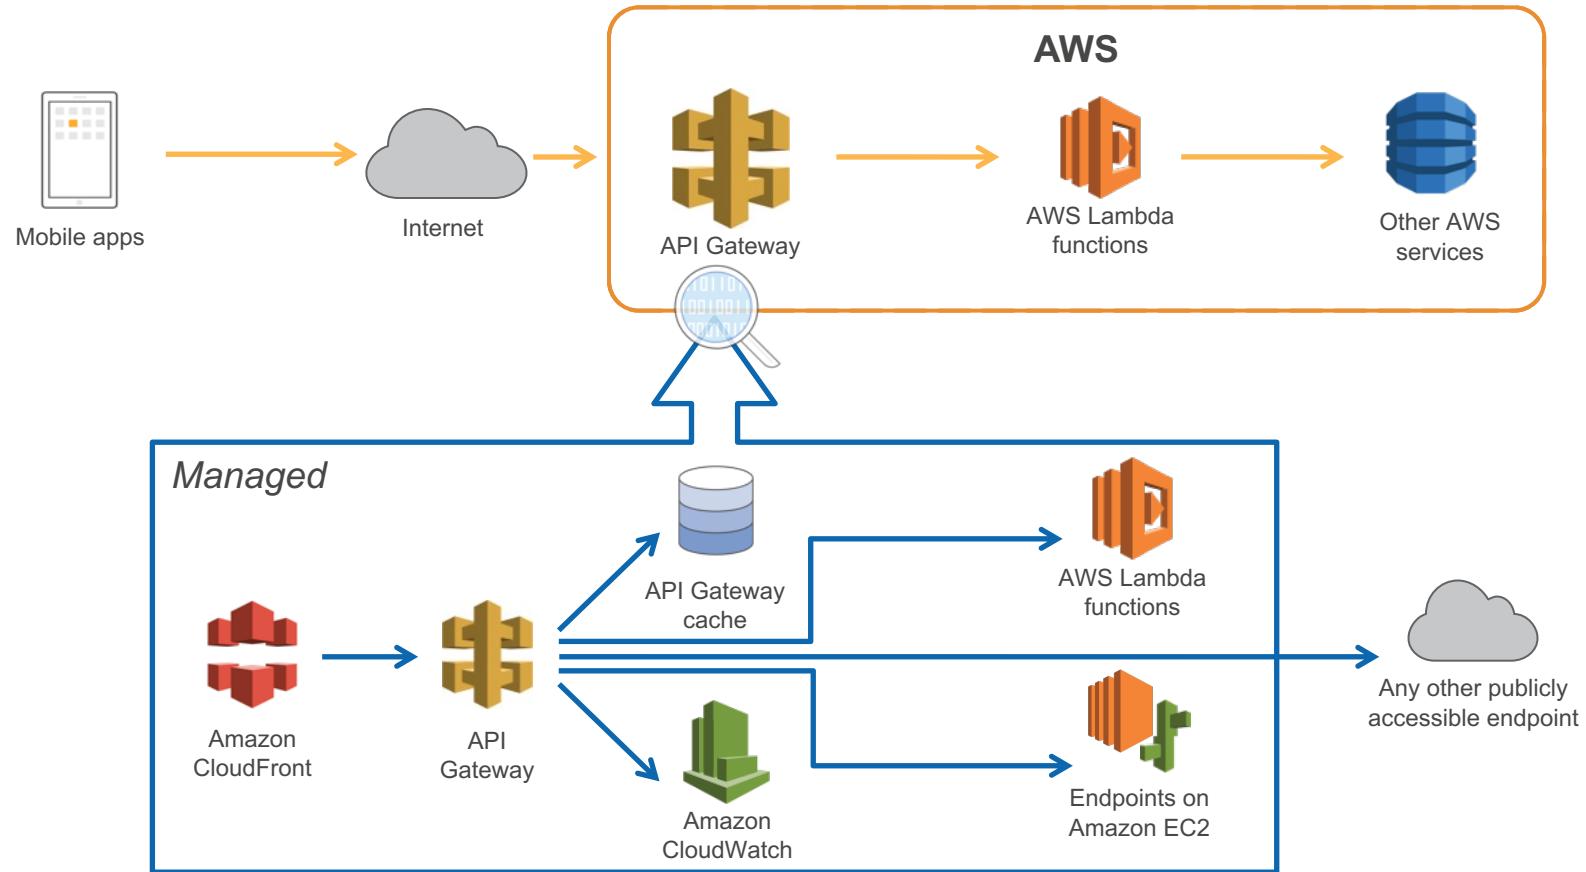


# Introducing Amazon API Gateway



- ✓ Managed cache to store API responses
- ✓ Reduced latency and DDoS protection through CloudFront
- ✓ SDK Generation for iOS, Android and JavaScript
- ✓ Swagger support
- ✓ Request / Response data transformation and API mocking

# A new, fully managed model



# Build, Deploy, Clone & Rollback

Build APIs with their resources, methods, and settings

Deploy APIs to a Stage

- Users can create as many Stages as they want, each with its own Throttling, Caching, Metering, and Logging configuration

Clone an existing API to create a new version

- Users can continue working on multiple versions of their APIs

Rollback to previous deployments

- We keep a history of customers' deployments so they can revert to a previous deployment

# API Configuration

You can create APIs

Define resources within an API

Define methods for a resource

- Methods are Resource + HTTP verb

Pet Store

/pets

/pets/{petId}

- GET
- POST
- PUT

# API Deployments

API Configuration can be deployed to a stage

Stages are different environments

For example:

- Dev (e.g. awsapigateway.com/dev)
- Beta (e.g. awsapigateway.com/beta)
- Prod (e.g. awsapigateway.com/prod)
- As many stages as you need

Pet Store

dev

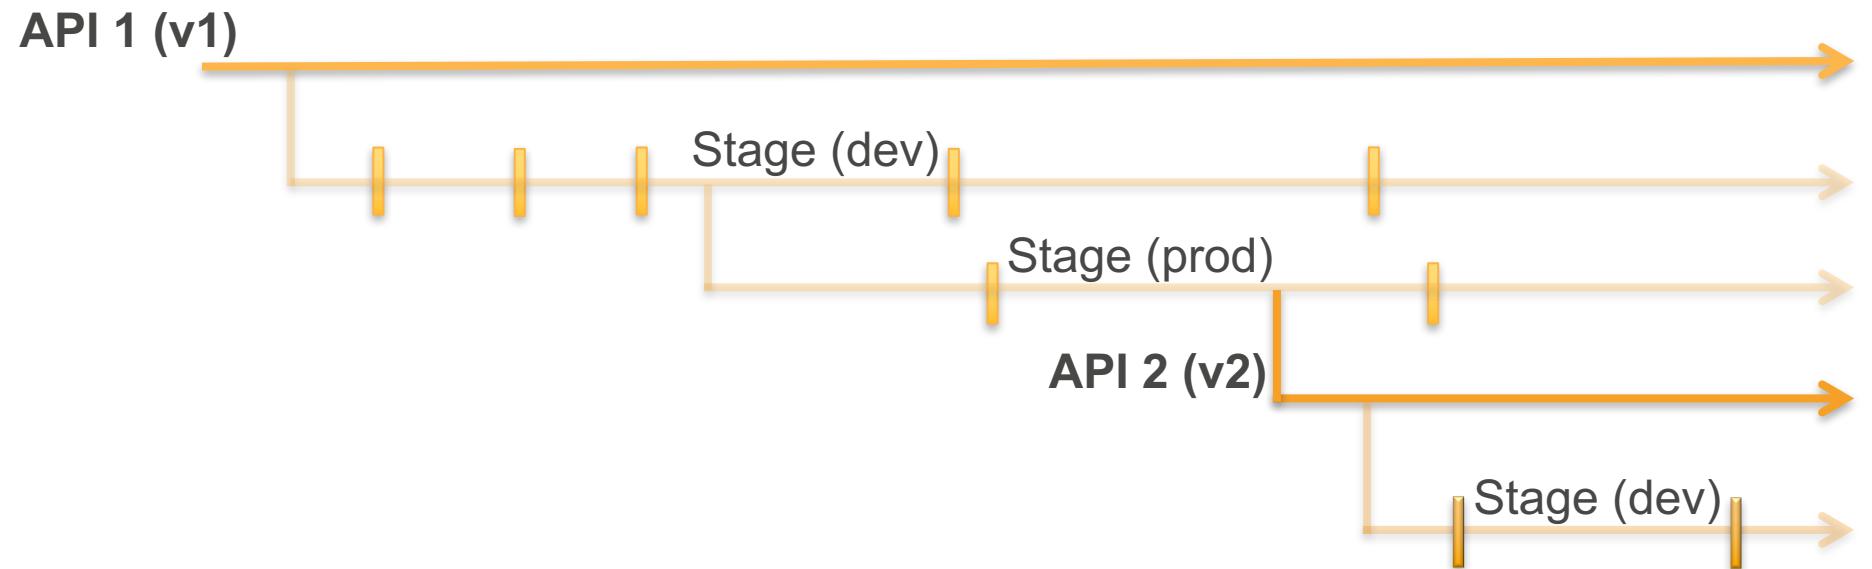
beta

gamma

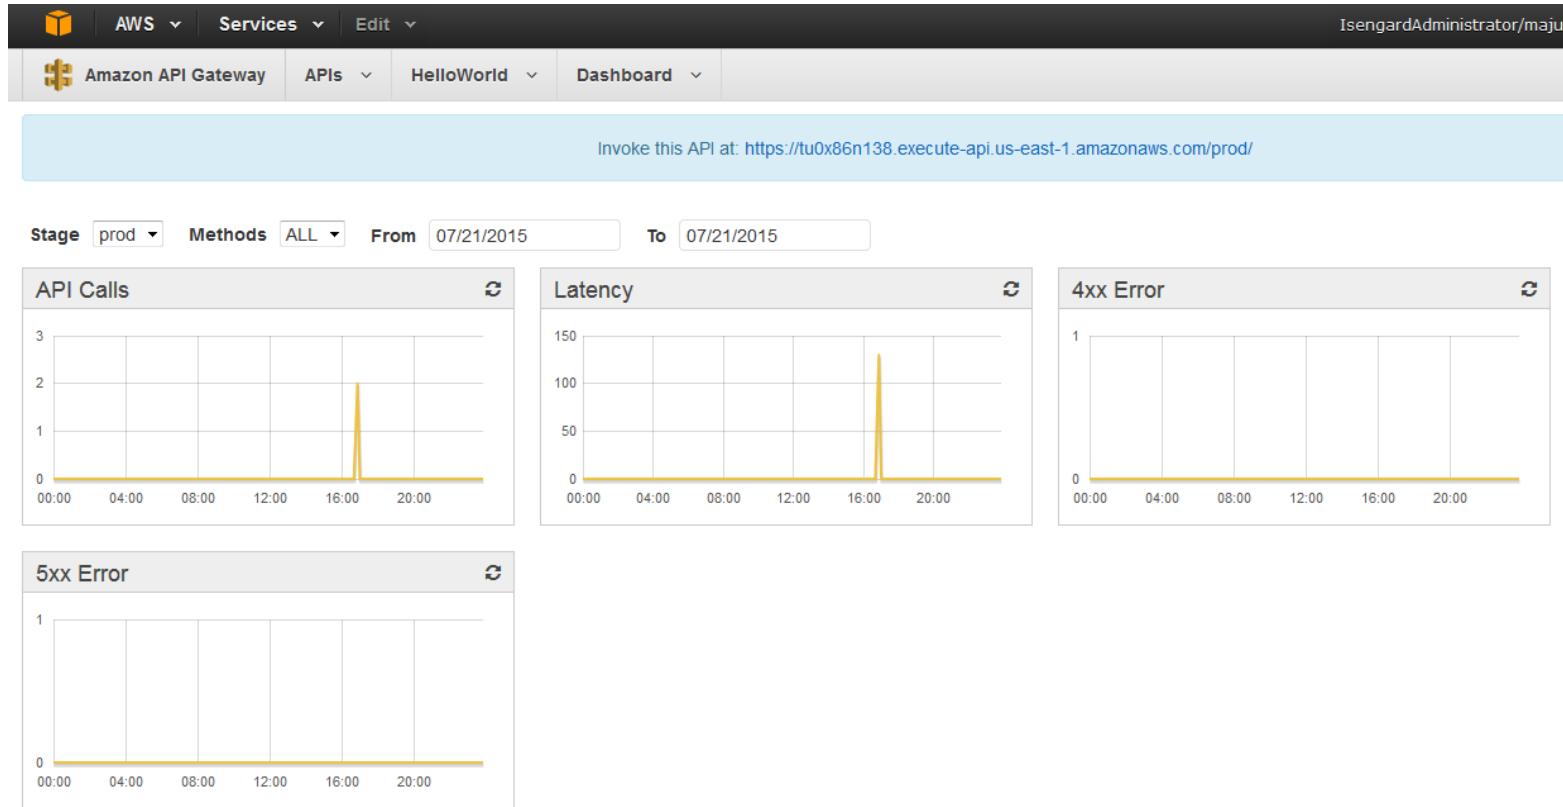
prod



# Manage Multiple Versions and Stages of your APIs



# API Monitoring



# API Throttling

Throttling helps you manage traffic to your backend

Throttle by developer-defined Requests/Sec limits

Requests over the limit are throttled

- HTTP 429 response

The generated SDKs retry throttled requests

# Caching of API Responses

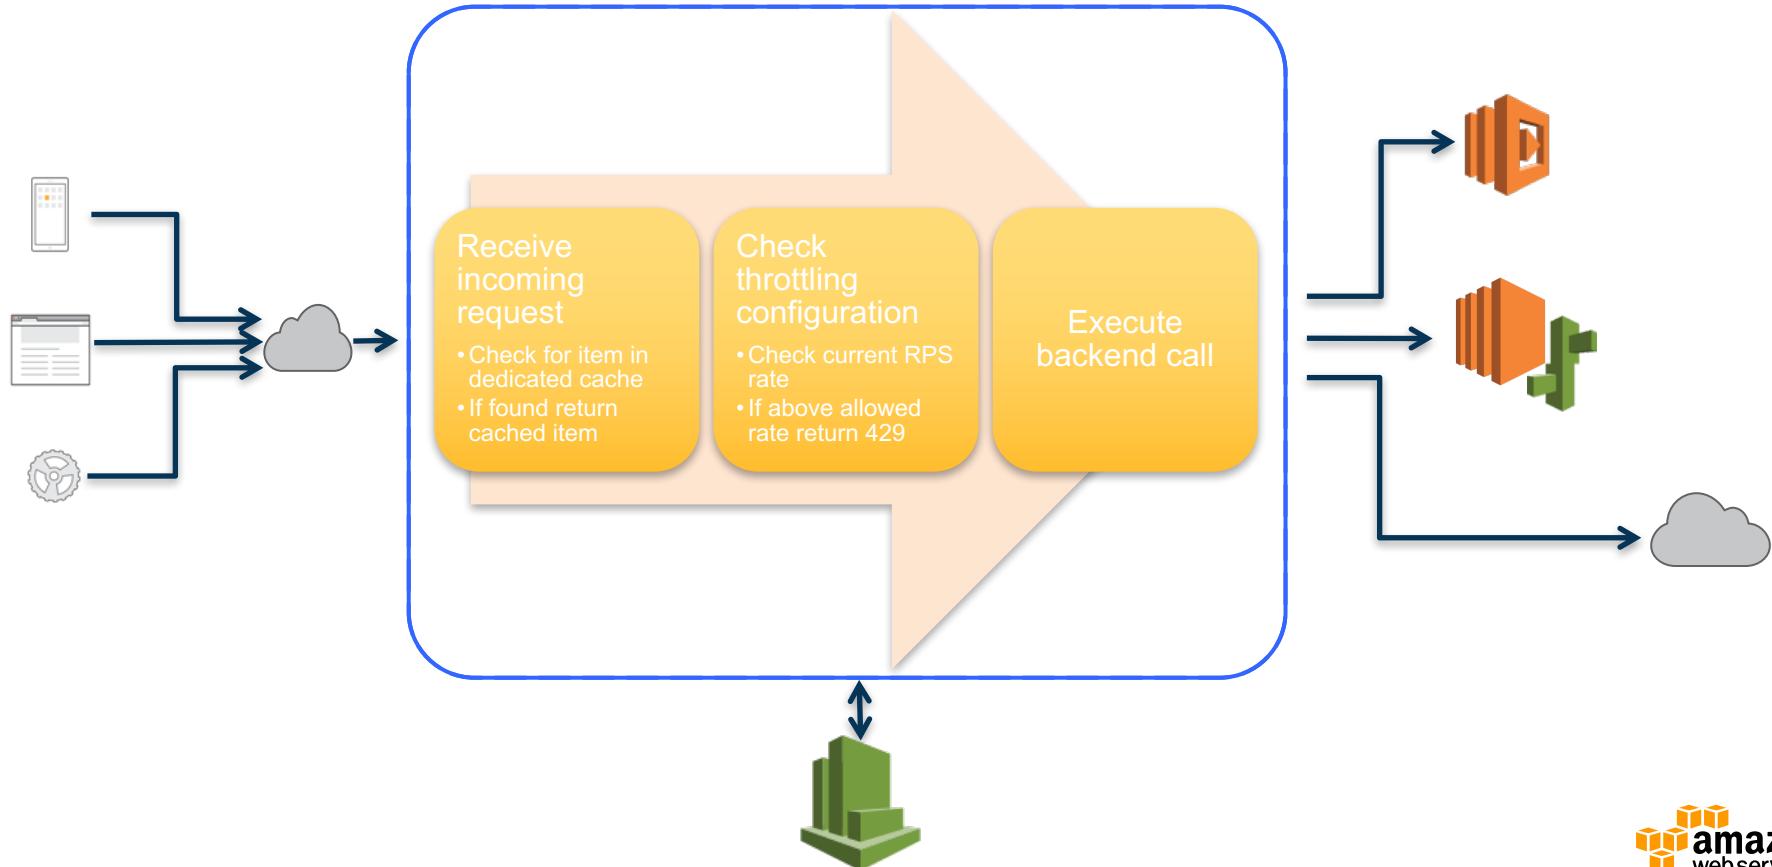
You can configure a cache key and the Time to Live (TTL) of the API response

Cached items are returned without calling the backend

A cache is dedicated to you, by stage

You can provision between 0.5GB to 237GB of cache

# Request processing workflow



# Identity and Access Management

# The IAM role defines access permissions

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "dynamodb:GetItem",  
        "dynamodb:PutItem",  
        "dynamodb:Scan",  
        "lambda:InvokeFunction",  
        "execute-api:invoke"  
      ],  
      "Resource": [  
        "arn:aws:dynamodb:us-east-1:xxxxxx:table/test_pets",  
        "arn:aws:lambda:us-east-1:xxxxxx:function:PetStore",  
        "arn:aws:execute-api:us-east-1:xxxx:API_ID/*/POST/pets"  
      ]  
    }  
  ]  
}
```

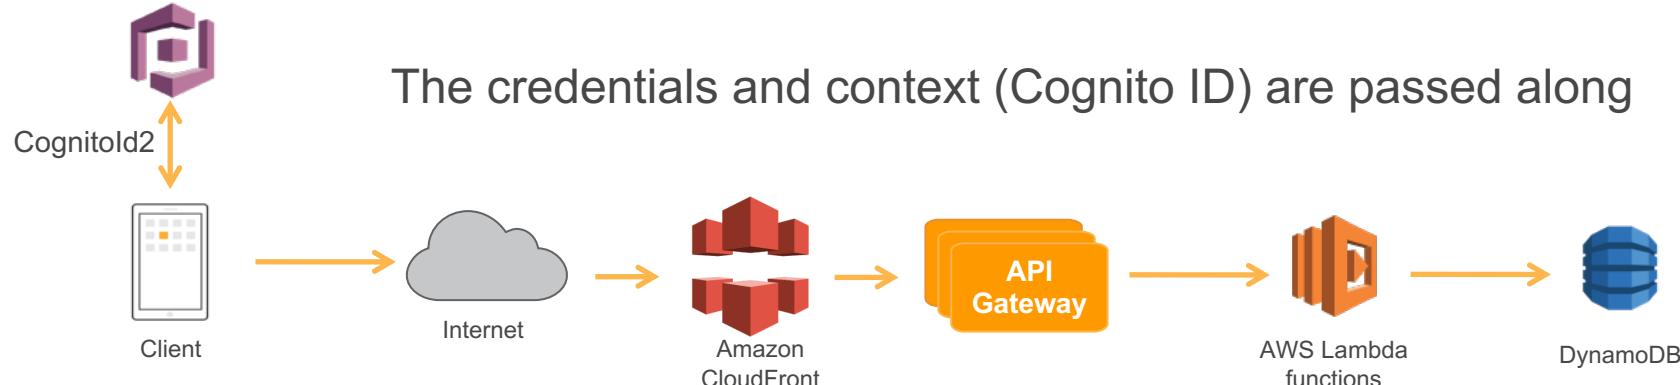
The role allows calls to:

- DynamoDB
- API Gateway
- Lambda

The role can access specific resources in these services

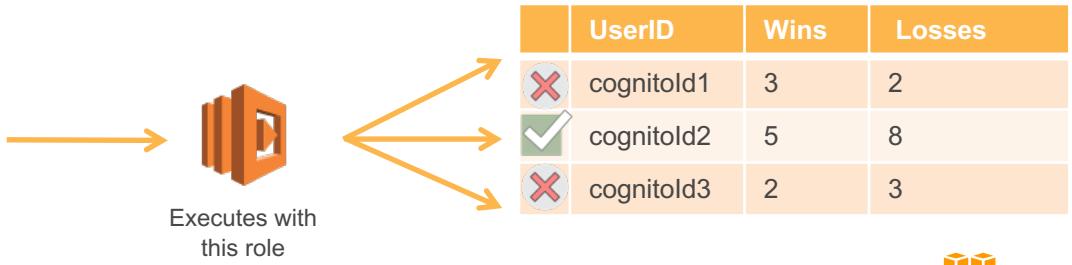


# One step further: Fine-grained access permissions



Both AWS Lambda & DynamoDB will follow the access policy

```
...  
    "Condition": {  
        "ForAllValues:StringEquals": {  
            "dynamodb:LeadingKeys": ["${cognito-  
identity.amazonaws.com:sub}"],  
            "dynamodb:Attributes": [  
                "UserId", "GameTitle", "Wins", "Losses",  
                "TopScore", "TopScoreDateTime"  
            ]  
        },  
        "StringEqualsIfExists": {  
            "dynamodb:Select": "SPECIFIC_ATTRIBUTES"  
        }  
    }  
...
```



# Authenticated flow in depth



*Learn more about fine-grained access permissions*

<http://amzn.to/1YkxcjR>

# API Gateway Authentication

IAM credentials per API method

AWS Sigv4 to sign and authorize API calls

Temporary credentials

- Via Amazon Cognito or AWS Security Token Service (STS)

```
POST https://iam.amazonaws.com/ HTTP/1.1
host: iam.amazonaws.com
Content-type: application/x-www-form-urlencoded; charset=utf-8
x-amz-date: 20110909T233600Z
```

```
Action=ListUsers&Version=2010-05-08
```



```
POST https://iam.amazonaws.com/ HTTP/1.1
Authorization: AWS4-HMAC-SHA256 Credential=AKIEXAMPLE/20110909/us-east-1/iam/aws4_requ
host: iam.amazonaws.com
```

```
Content-type: application/x-www-form-urlencoded; charset=utf-8
x-amz-date: 20110909T233600Z
```

```
Action=ListUsers&Version=2010-05-08
```

# Authentication via Custom Header

OAuth or other authorization mechanisms through custom headers

Simply configure your API methods to forward the custom headers to your backend

# Benefits of using AWS auth & IAM

- Separation of concerns – our **authorization strategy is delegated** to a dedicated service
- We have **centralized access management** to a single set of policies
- **Roles and credentials can be disabled** with a single API call

# API Keys to Meter Developer Usage

Create API Keys

Set access permissions at the API/Stage level

Meter usage of the API Keys through CloudWatch Logs

# API Keys to Authorize Access



The name “Key” implies security – there is no security in baking text in an App’s code



API Keys should be used purely to meter app/developer usage



API Keys should be used alongside a stronger authorization mechanism

# Amazon API Gateway Best Practices

1. Use mock integrations
2. Combine with Amazon Cognito for managed end user-based access control.
3. Use stage variables (inject API config values into Lambda functions for logging, behavior).
4. Use request/response mapping templates everywhere within reason, not passthrough.
5. Take ownership of HTTP response codes.
6. Use Swagger import/export for cross-account sharing.



# API Gateway Demo

# AWS Lambda

# AWS Lambda

λ

Announced at re:Invent 2014

- Deploy **pure functions** in Java, Python, Node.js and C#
- Just **code**, without the infrastructure drama
- Built-in **scalability** and **high availability**
- **Integrated** with many AWS services
- **Pay as you go**
  - Combination of execution time (100ms slots) & memory
  - Starts at \$0.000000208 per 100ms
  - Free tier available: first 1 million requests per month are free



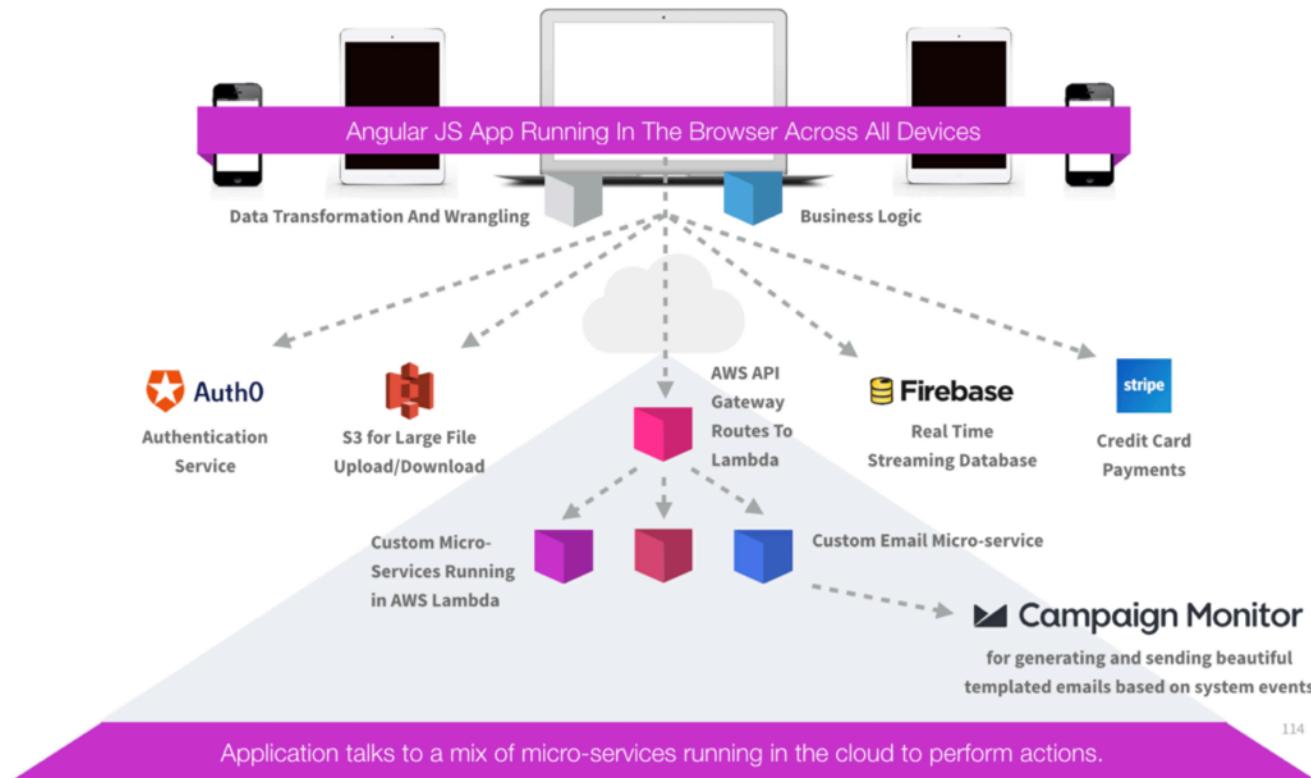
# What can you do with AWS Lambda?



- Grow ‘**connective tissue**’ in your AWS infrastructure
  - Example: <http://www.slideshare.net/JulienSIMON5/building-a-serverless-pipeline>
- Build **event-driven** applications
- Build **APIs** together with Amazon API Gateway
  - RESTful APIs
  - Resources, methods
  - Stages



# A Cloud Guru: 100% Serverless



# AWS Lambda Programming Model



## Bring your own code

- Node.js, Java, Python, C#
- Bring your own libraries (even native ones)



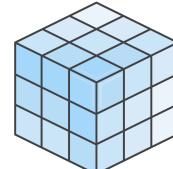
## Simple resource model

- Select power rating from 128 MB to 1.5 GB
- CPU and network allocated proportionately
- Reports actual usage



## Programming model

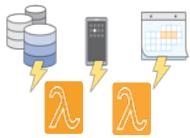
- AWS SDK built in (Python and Node.js)
- Lambda is the “webserver”
- Use processes, threads, /tmp, sockets normally



## Stateless

- Persist data using Amazon DynamoDB, S3, or ElastiCache
- No affinity to infrastructure (can't “log in to the box”)

# Using AWS Lambda



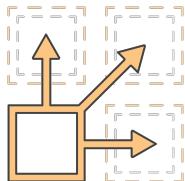
## Flexible use

- Call or send events
- Integrated with other AWS services
- Build whole serverless ecosystems



## Flexible authorization

- Securely grant access to resources, including VPCs
- Fine-grained control over who can call your functions



## Authoring functions

- Author directly using the console WYSIWYG editor
- Package code as a .zip and upload to Lambda or S3
- Plugins for Eclipse and Visual Studio
- Command line tools



## Monitoring and logging

- Built-in metrics for requests, errors, latency, and throttles
- Built-in logs in Amazon CloudWatch Logs

# AWS Lambda Best Practices

1. Limit your function size – especially for Java (starting the JVM takes time).
2. Don't assume function container reuse – but take advantage of it when it does occur.
3. Don't forget about disk (500 MB /tmp directory provided to each function).
4. Use function aliases for release.
5. Use the included logger (include details from service-provided context).
6. Create custom metrics (operations-centric, and business-centric).



# Amazon DynamoDB

# Amazon DynamoDB

Run your business, not your database



# DynamoDB Benefits



-  Fully managed
-  Fast, consistent performance
-  Highly scalable
-  Flexible
-  Event-driven programming
-  Fine-grained access control

# Fully managed service = automated operations

- App optimization
- Scaling
- High availability
- Database backups
- DB s/w patches
- DB s/w installs
- OS patches
- OS installation
- Server maintenance
- Rack & stack
- Power, HVAC, net

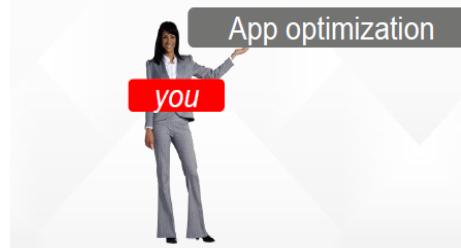


DB hosted on premise



- App optimization

*you*



DynamoDB

- Scaling
- High availability
- Database backups
- DB s/w patches
- DB s/w installs
- OS patches
- OS installation
- Server maintenance
- Rack & stack
- Power, HVAC, net



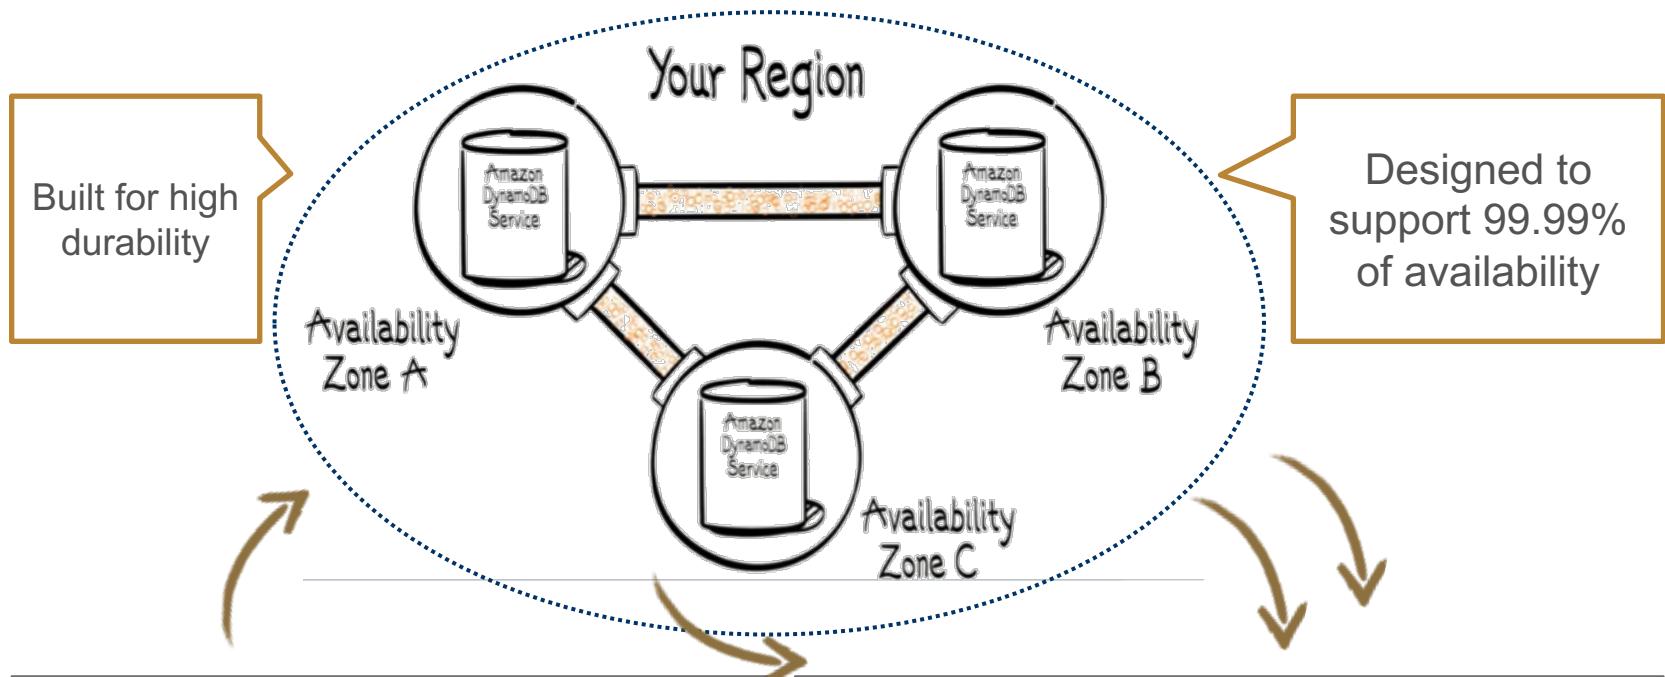
# Consistently low latency at scale



PREDICTABLE  
PERFORMANCE!



# High availability and durability



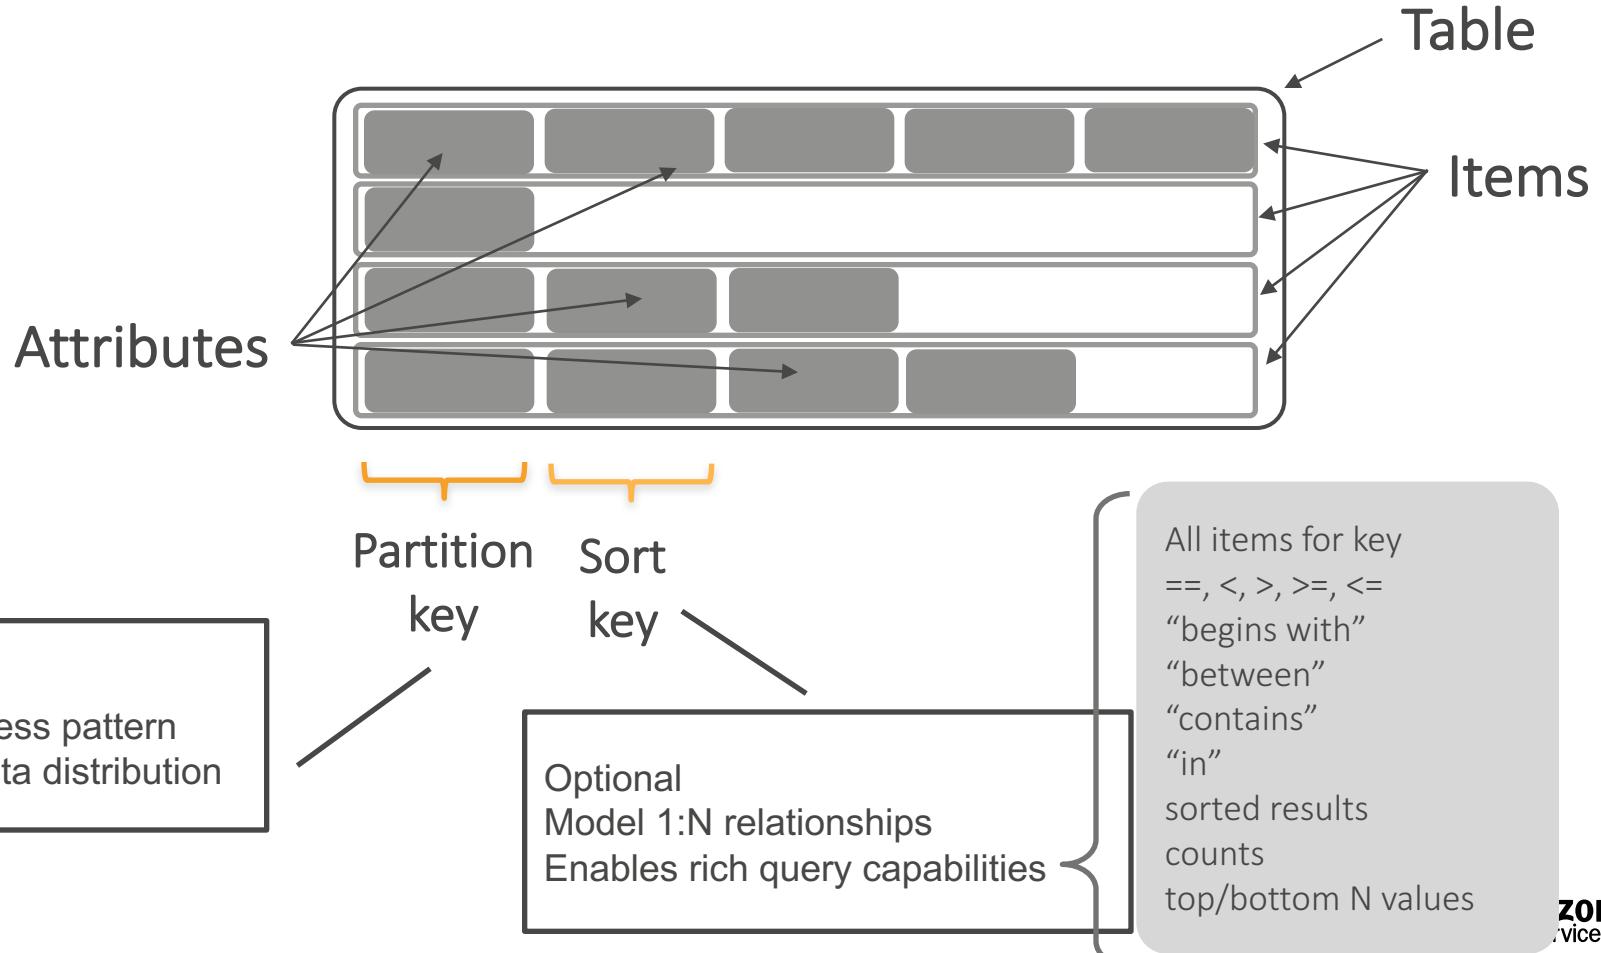
## WRITES

Replicated continuously to 3 AZs  
Persisted to disk (custom SSD)

## READS

Strongly or eventually consistent  
No latency trade-off

# DynamoDB table structure



# Advanced topics in DynamoDB

- Data modeling
- Indexes - query a table using alternate attributes
  - Local Secondary Indexes
  - Global Secondary Indexes
- Understanding Partitions
  - # of partitions depend on table throughput and size
- Design patterns and best practices

# Scaling

## Throughput

- Provision any amount of throughput to a table

## Size

- Add any number of items to a table
  - Maximum item size is 400 KB
  - LSIs limit the number of range keys due to 10 GB limit

Scaling is achieved through partitioning

# Throughput

Provisioned at the table level

- Write capacity units (WCUs) are measured in 1 KB per second
- Read capacity units (RCUs) are measured in 4 KB per second
  - RCUs measure strictly consistent reads
  - Eventually consistent reads cost 1/2 of consistent reads

Read and write throughput limits are independent



RCU



WCU

# Simplifying Development

# Typical development workflow

1. Write and deploy a Lambda function
2. Create and deploy a REST API with API Gateway
3. Connect the API to the Lambda function
4. Invoke the API
5. Test, debug and repeat ;)

# The Serverless framework

formerly known as JAWS: Just AWS Without Servers



- Announced at **re:Invent 2015** by Austen Collins and Ryan Pendergast
- Supports **Node.js**, as well as **Python** and **Java** (with restrictions)
- Auto-deploys and runs **Lambda functions** **locally** or **remotely**
- Auto-deploys your **Lambda event sources**: API Gateway, S3, DynamoDB
- Creates all required infrastructure with **CloudFormation**
- Simple configuration in **YML**

<http://github.com/serverless/serverless>

<https://serverless.com>

AWS re:Invent 2015 | (DVO209) [https://www.youtube.com/watch?v=D\\_U6luQ6I90](https://www.youtube.com/watch?v=D_U6luQ6I90) <https://vimeo.com/141132756>



# Serverless: “Hello World” API

```
$ serverless create
```

*Edit handler.js, serverless.yml and event.json*

```
$ serverless deploy [--stage stage_name]
```

```
$ serverless invoke [--local] --function function_name
```

```
$ serverless info
```

```
$ http $URL
```

# Gordon

Released in Oct'15 by Jorge Batista

Supports **Python, Javascript, Golang, Java, Scala, Kotlin** (including in the same project)

Auto-deploys and runs **Lambda functions**, locally or remotely

Auto-deploys your **Lambda event sources**: API Gateway, CloudWatch Events, DynamoDB Streams, Kinesis Streams, S3

Creates all required infrastructure with **CloudFormation**

Simple configuration in **YML**

# Gordon: “Hello World” API

```
$ gordon startproject helloworld
```

```
$ gordon startapp helloapp
```

*Write hellofunc() function*

```
$ gordon build
```

```
$ echo '{"name": "Julien"}' | gordon run helloapp.hellofunc
```

```
$ gordon apply [--stage stage_name]
```

```
$ http post $URL name=Julien
```



# AWS Chalice

Think of it as a serverless framework for Flask apps

- Released in Jul'16, still in **beta**
- Just add **your Python code**
  - Deploy with a **single call** and **zero config**
  - The API is created **automatically**, the IAM policy is **auto-generated**
- Run APIs **locally** on port 8000 (similar to Flask)
- **Fast & lightweight** framework
  - 100% *boto3* calls (AWS SDK for Python) → fast
  - No integration with CloudFormation → no creation of event sources



# AWS Chalice: “Hello World” API

```
$ chalice new-project helloworld
```

*Write your function in app.py*

```
$ chalice local
```

```
$ chalice deploy
```

```
$ export URL=`chalice url`
```

```
$ http $URL
```

```
$ http put $URL/hello/tobias
```

```
$ chalice logs [ --include-lambda-messages ]
```



# AWS Chalice: PUT/GET in S3 bucket

```
$ chalice new-project s3test
```

*Write your function in app.py*

```
$ chalice local
```

```
$ http put http://localhost:8000/objects/doc.json value1=5 value2=8
```

```
$ http get http://localhost:8000/objects/doc.json
```

```
$ chalice deploy [stage_name]
```

```
$ export URL=`chalice url`
```

```
$ http put $URL/objects/doc.json value1=5 value2=8
```

```
$ http get $URL/objects/doc.json
```



# Demo Chalice

# Summing things up

## Serverless

The most popular serverless framework

Built with and for Node.js.  
Python and Java: YMMV

Rich features, many event sources

Not a web framework

## Gordon

Great challenger!

Node.js, Python, Java,  
Scala, Golang

Comparable to Serverless feature-wise

Not a web framework

## Chalice

AWS project, in beta

Python only

Does only one thing, but does it great

Dead simple, zero config

Flask web framework

# More Lambda frameworks

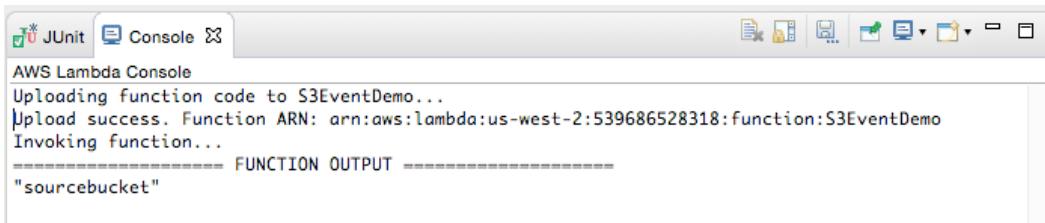
- **Kappa** <https://github.com/garnaat/kappa>
  - Released Dec'14 by Mitch Garnaat, author of boto and the AWS CLI (still maintained?)
  - Python only, multiple event sources
- **Apex** <https://github.com/apex/apex>
  - Released in Dec'15 by TJ Holowaychuk
  - Python, Javascript, Java, Golang
  - Terraform integration to manage infrastructure for event sources
- **Zappa** <https://github.com/Miserlou/Zappa>
  - Released in Feb'16 by Rich Jones
  - Python web applications on AWS Lambda + API Gateway
- **Docker-lambda** <https://github.com/lambci/docker-lambda>
  - Released in May'16 by Michael Hart
  - Run functions in Docker images that “replicate” the live Lambda environment



# 2 Java tools for AWS Lambda

## Eclipse plug-in

- Code test and deploy Lambdas from Eclipse
- Run your functions locally and remotely
- Test with local events and JUnit4
- Deploy standalone functions, or with the AWS Serverless Application Model (Dec'16)



<https://java.awsblog.com/post/TxWZES6J1RSQ2Z/Testing-Lambda-functions-using-the-AWS-Toolkit-for-Eclipse>

<https://aws.amazon.com/blogs/developer/aws-toolkit-for-eclipse-serverless-application>

<https://github.com/awslabs/aws-serverless-java-container>

## Serverless Java Container

- Run Java RESTful APIs as-is
- Default implementation of the Java servlet  
    HttpServletResponse  
    HttpServletRequest
- Support for Java frameworks such as Jersey or Spark



# Lambda Java Resources

- [AWS Lambda Getting Started](#)
- [Lambda Java Programming Model](#)
- [AWS Java Developer Blog](#)
- [AWS Serverless Java Container](#)
- [POJO Modeling with DynamoDBMapper](#)
- [Lambada Framework](#)

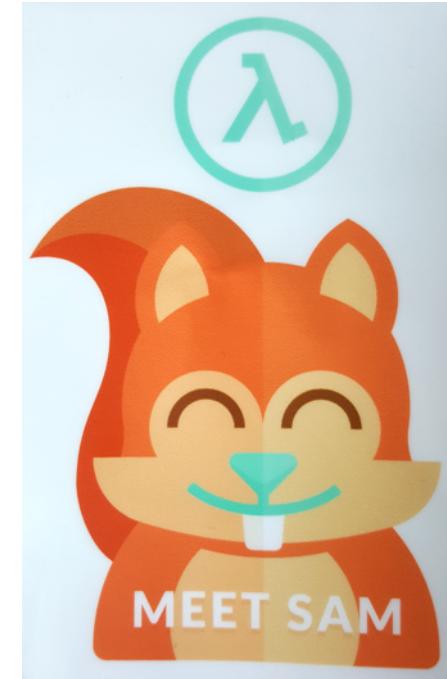
# Best Practices

1. Use strategic, consumable naming conventions (Lambda function names, IAM roles, API names, API stage names, etc.).
2. Use naming conventions and versioning to create automation.
3. Externalize authorization to IAM roles whenever possible.
4. Least privilege and separate IAM roles.
5. Externalize configuration – DynamoDB is great for this.
6. Be aware of service throttling, engage AWS support if so.

# Simplifying Deployment

# AWS Serverless Application Model (SAM)

- CloudFormation extension released in Nov'16 to bundle Lambda functions, APIs & events
- 3 new CloudFormation resource types
  - AWS::Serverless::Function
  - AWS::Serverless::Api
  - AWS::Serverless::SimpleTable (DynamoDB)
- 2 new CloudFormation CLI commands
  - ‘aws cloudformation package’
  - ‘aws cloudformation deploy’
- Integration with CodeBuild and CodePipeline for CI/CD
- Expect SAM to be integrated in most / all frameworks



```
AWSTemplateFormatVersion: '2010-09-09'
Transform: AWS::Serverless-2016-10-31
Description: Get items from a DynamoDB table.

Resources:
  GetFunction:
    Type: AWS::Serverless::Function
    Properties:
      Handler: index.get
      Runtime: nodejs4.3
      Policies: AmazonDynamoDBReadOnlyAccess
    Environment:
      Variables:
        TABLE_NAME: !Ref Table
  Events:
    GetResource:
      Type: Api
      Properties:
        Path: /resource/{resourceId}
        Method: get
  Table:
    Type: AWS::Serverless::SimpleTable
```

Sample SAM template for:

- Lambda function
- HTTP GET API
- DynamoDB table



# SAM: Open Specification

A common language to  
describe the content of a  
serverless application  
*across the ecosystem.*

Apache 2.0 licensed  
GitHub project

## AWS Serverless Application Model (SAM)

Version 2016-10-31

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC 2119](#).

The AWS Serverless Application Model (SAM) is licensed under [The Apache License, Version 2.0](#).

### Introduction

AWS SAM is a model used to define serverless applications on AWS.

Serverless applications are applications composed of functions triggered by events. A typical serverless application consists of one or more AWS Lambda functions triggered by events such as object uploads to [Amazon S3](#), API calls to [Amazon API Gateway](#), and scheduled events. Those functions can stand alone or leverage other resources such as [Amazon DynamoDB](#), [Amazon SNS](#), and [Amazon SQS](#). The most basic serverless application is simply a function.



# Demo AWS Toolkit for Eclipse

# New Lambda videos from re:Invent 2016

AWS re:Invent 2016: What's New with AWS Lambda (SVR202)

<https://www.youtube.com/watch?v=CwxWhyGteNc>

AWS re:Invent 2016: Serverless Apps with AWS Step Functions (SVR201)

<https://www.youtube.com/watch?v=75MRve4nv8s>

AWS re:Invent 2016: Real-time Data Processing Using AWS Lambda (SVR301)

<https://www.youtube.com/watch?v=VFLKOy4GKXQ>

AWS re:Invent 2016: Serverless Architectural Patterns and Best Practices (ARC402)

<https://www.youtube.com/watch?v=b7UMoc1iUYw>

AWS re:Invent 2016: Bringing AWS Lambda to the Edge (CTD206)

<https://www.youtube.com/watch?v=j26novaqF6M>

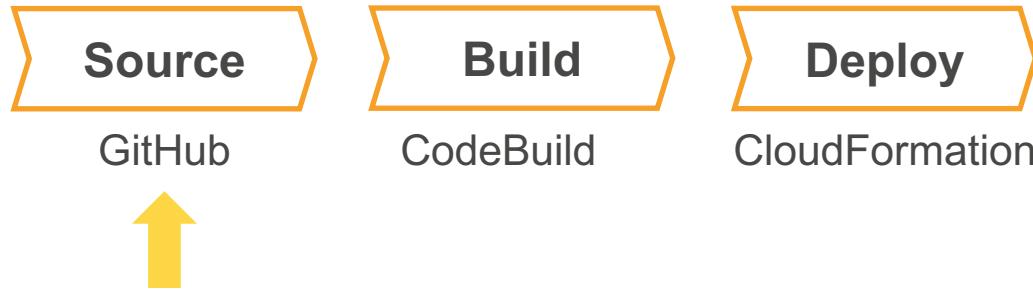
AWS re:Invent 2016: Ubiquitous Computing with Greengrass (IOT201)

<https://www.youtube.com/watch?v=XQQjX8GTEko>



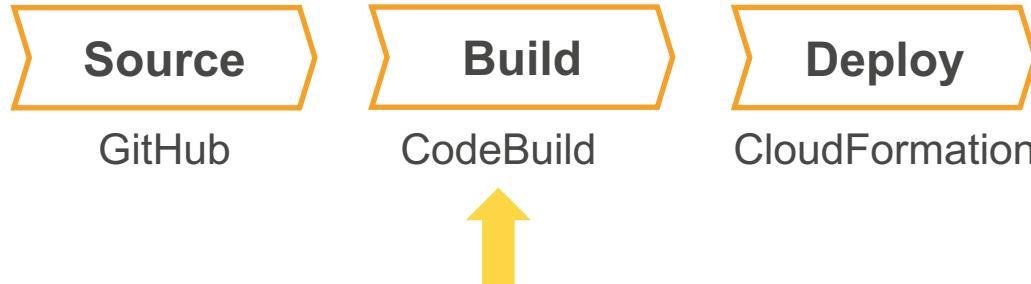
# Putting it all together

# Serverless CI/CD pipeline



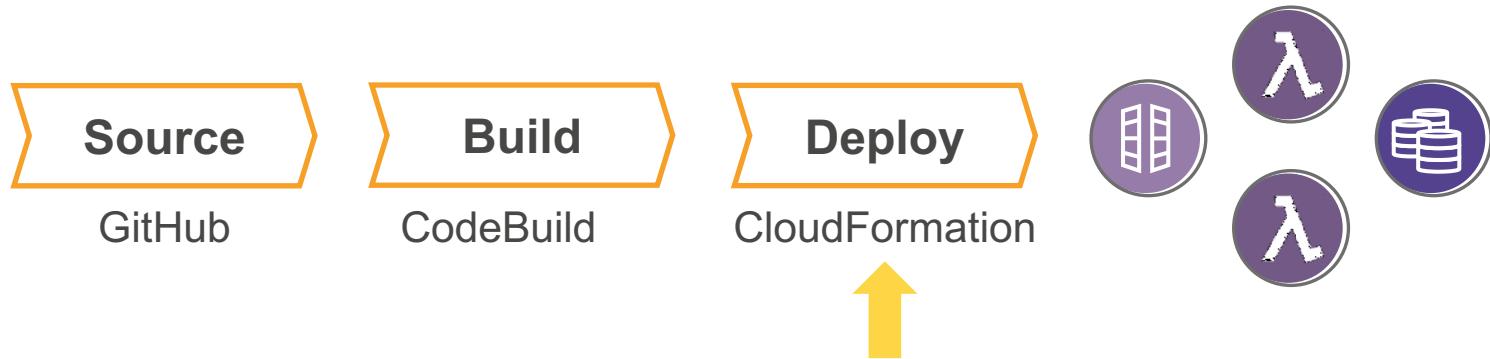
- Pull source directly from GitHub or CodeCommit using CodePipeline

# Serverless CI/CD pipeline



- Pull source directly from GitHub or AWS CodeCommit using AWS CodePipeline
- Build and package serverless apps with AWS CodeBuild
  - npm, pip, Java compilation, BYO Docker...

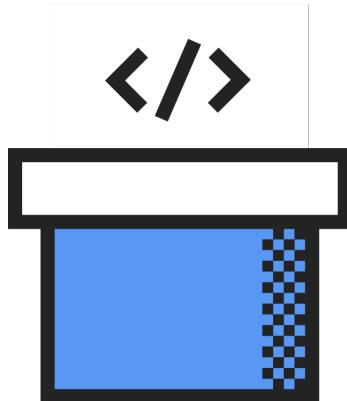
# Serverless CI/CD pipeline



- Pull source directly from GitHub or AWS CodeCommit using AWS CodePipeline
- Build and package serverless apps with AWS CodeBuild
- Deploy your completed Lambda app with AWS CloudFormation

# AWS CodePipeline

Continuous delivery service for fast and reliable application updates



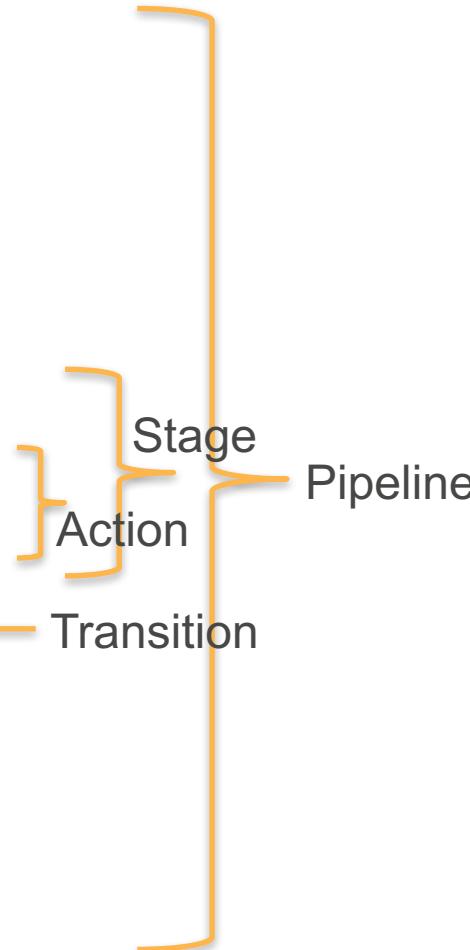
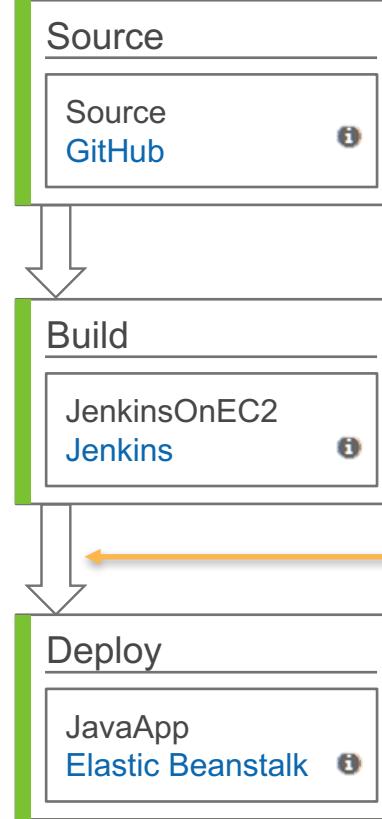
Model and visualize your software release process

Builds, tests, and deploys your code every time there is a code change

Integrates with third-party tools and AWS

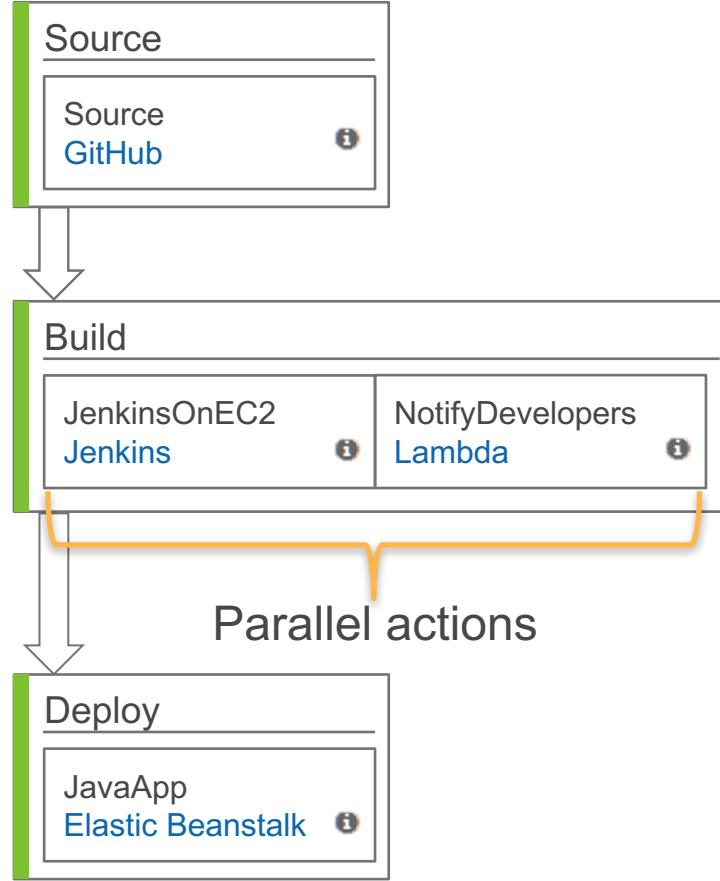


MyApplication



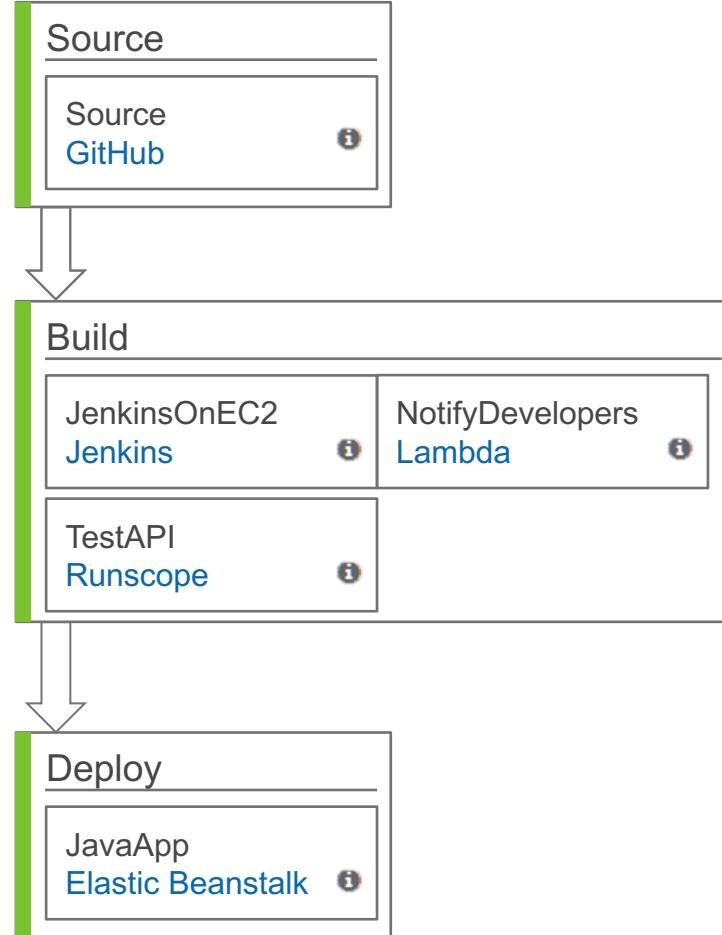


## MyApplication





## MyApplication

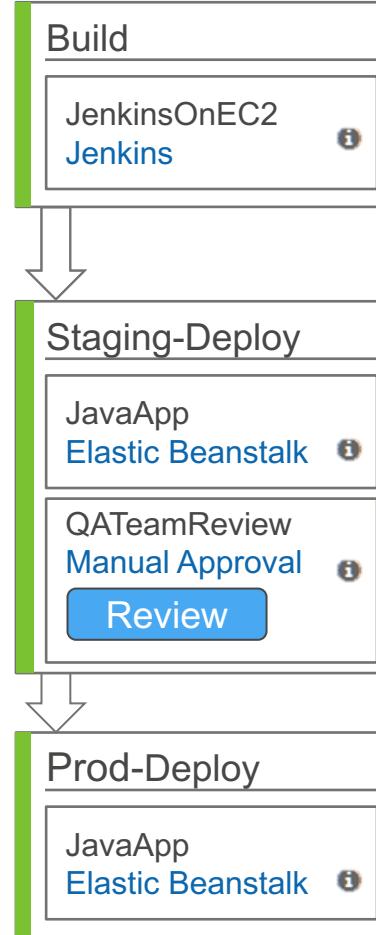


Sequential actions





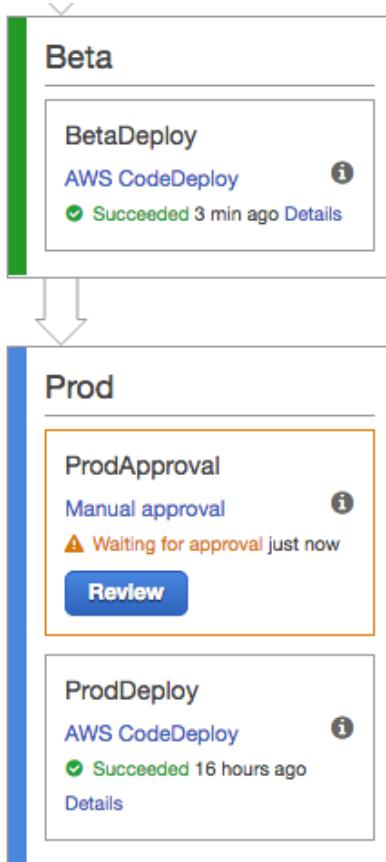
## MyApplication



Manual approvals



# Manual approvals

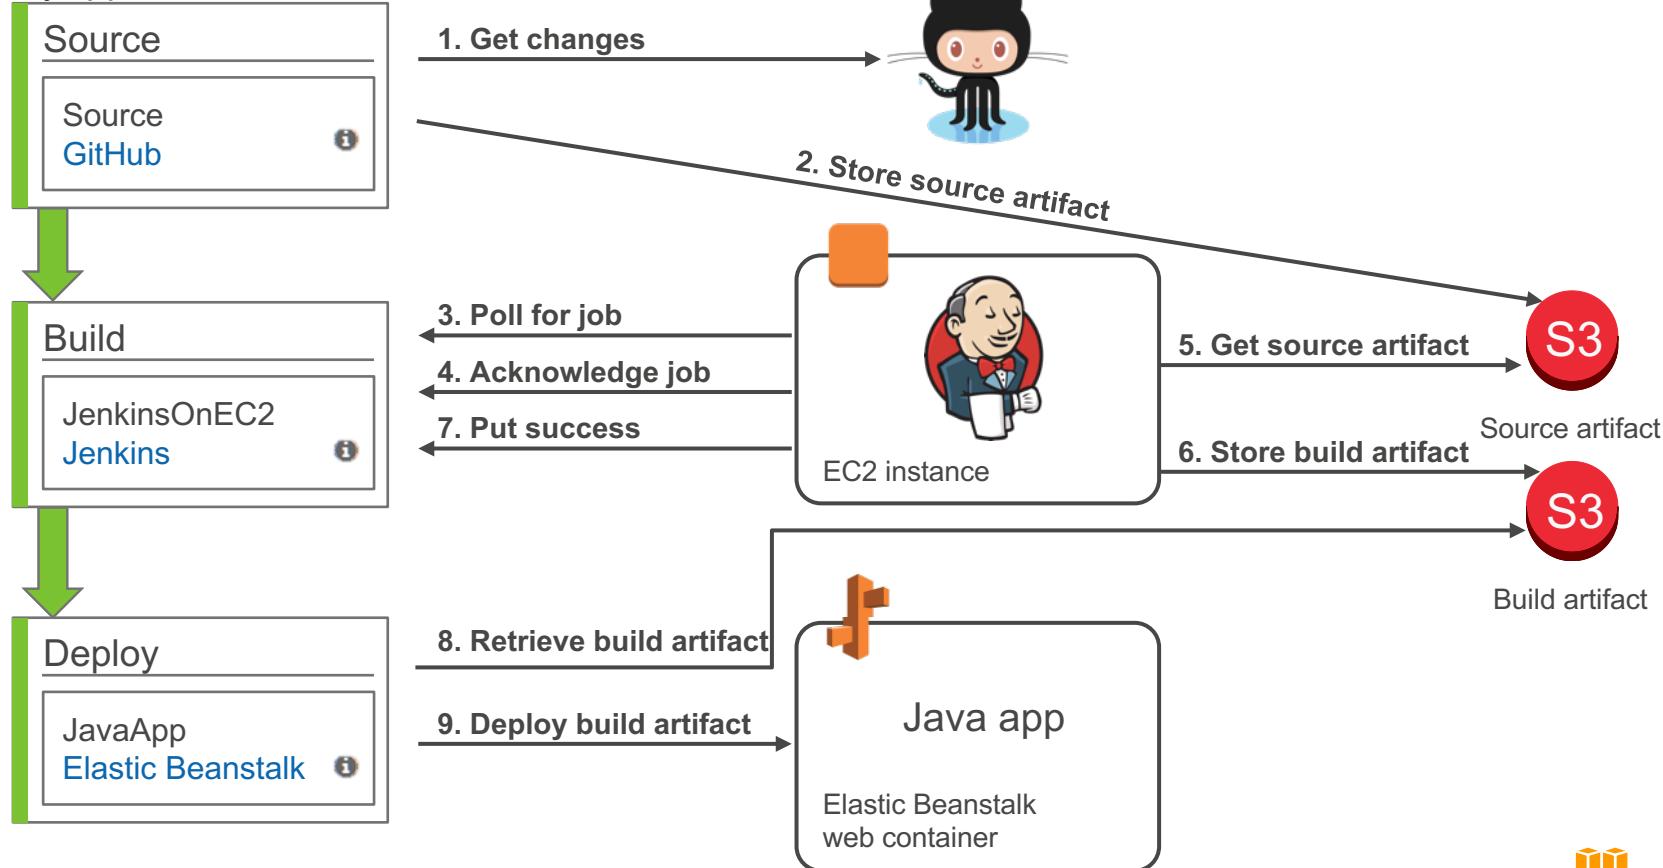


You can add a manual approval at the point where you want the pipeline to stop running until someone approves or rejects the revision in progress.

- Pipeline will stop executing when it has reached the point at which you set the approval action
- Pipeline execution resumes only when the action has been approved
- Approval action managed with AWS Identity and Access Management (IAM) permissions
- Notify approvers in several ways including email, SMS, webhooks, and more
- Useful for manual QA actions or as part of “Canary” deploy models



## MyApplication



# We have a strong partner list, and it's growing

Source

Build

Test

Deploy

**GitHub**

  
CloudBees®

 Apica

 XebiaLabs  
Deliver Faster

 Jenkins

  
Solano Labs

 TeamCity  
\*beta

 BlazeMeter

 Ghost Inspector

---

 HPE StormRunner

 Runscope

 Amazon  
Web Services

# AWS service integrations

Source

Invoke Logic

Deploy

Amazon S3

AWS Lambda

AWS CodeDeploy

AWS CodeCommit

AWS Elastic Beanstalk

AWS OpsWorks

AWS CloudFormation



# Demo Jersey Serverless Application and CI/CD Pipeline

<https://github.com/sctobbor/aws-serverless-jersey-example>



# Additional Resources

- [SSO to AWS with AD FS and SAML](#)
- [AWS WAF Security Automations](#)
- [Setting up DynamoDB Local](#)

